



Conception Avancée de Bases de Données

pg_stats



(C) Diego Marzou



D'après C.J DATE

DDL : langage de définition des données; DML : langage de manipulation des données

PostgreSQL's statistics collector



- PostgreSQL's statistics collector is a subsystem that supports collection and reporting of information about server activity.
- The collector can count accesses to tables and indexes in both disk-block and individual-row terms.

Postgresql documentation



- <http://www.postgresql.org/docs/9.2/static/view-pg-stats.html>
- The view `pg_stats` provides access to the information stored in the `pg_statistic` catalog.
- This view allows access only to rows of `pg_statistic` that correspond to tables the user has permission to read, and therefore it is safe to allow public read access to this view.

pg_statistic



- The catalog `pg_statistic` stores statistical data about the contents of the database.
- Entries are created by `ANALYZE` and subsequently used by the query planner.
- There is one entry for each table column that has been analyzed.
- Note that all the statistical data is inherently approximate, even assuming that it is up-to-date.

pgstat_stat_interval



- Each individual server process transmits new block and row access counts to the collector just before going idle;
- So a query or transaction still in progress does not affect the displayed totals.
- Also, the collector itself emits a new report at most once per `pgstat_stat_interval` milliseconds (500 by default). So the displayed information lags behind actual activity.
- Current-query information is reported to the collector immediately, but is still subject to the `pgstat_stat_interval` delay before it becomes visible.

<http://www.postgresql.org/docs/8.0/interactive/monitoring-stats.html>

pg_stats



```
emmanuelfuchs — psql — 84x20
cavbd=> \d pg_stats
        View "pg_catalog.pg_stats"

```

Column	Type	Modifiers
schemaname	name	
tablename	name	
attname	name	
inherited	boolean	
null_frac	real	
avg_width	integer	
n_distinct	real	
most_common_vals	anyarray	
most_common_freqs	real[]	
histogram_bounds	anyarray	
correlation	real	
most_common_elems	anyarray	
most_common_elem_freqs	real[]	
elem_count_histogram	real[]	


```
cavbd=> █
```

pg_stats view select



```
SELECT schemaname, tablename, attname, inherited, null_frac, avg_width,  
       n_distinct, most_common_vals, most_common_freqs, histogram_bounds,  
       correlation, most_common_elems, most_common_elem_freqs,  
       elem_count_histogram
```

```
FROM pg_stats;
```

Query - postgres sur postgres@localhost : 5432 *

Éditeur SQL Constructeur graphique de requêtes

Requêtes précédentes Supprimer Tout supprimer


```
SELECT tablename, atname, avg_width,  
       n_distinct, most_common_vals, most_common_freqs, histogram_bounds,  
       correlation  
  
FROM pg_stats;
```

Panneau sortie

Sortie de données Expliquer (Explain) Messages Historique

pg_class Hash Hash Join Hash Join pg_attribute

OK. Unix Ligne 4, Col 8, Caract. 141 11 lignes. 17 ms



Query - postgres sur postgres@localhost : 5432

postgres sur postgres@localhost : 5432

Éditeur SQL Constructeur graphique de requêtes

Requêtes précédentes Supprimer Tout supprimer

```
SELECT schemaname, tablename, attname, inherited, null_frac, avg_width,
       n_distinct, most_common_vals, most_common_freqs, histogram_bounds,
       correlation, most_common_elems, most_common_elem_freqs, elem_count_histogram
FROM pg_stats;
```

Panneau sortie

Sortie de données Expliquer (Explain) Messages Historique

pg_class Hash pg_statistic Hash Join pg_attribute Hash Join Nested Loop Left Join pg_namespace_oid_index

OK. Unix Ligne 1, Col 1, Caract. 1 26 lignes. 28 ms



Query - postgres sur postgres@localhost : 5432 *

postgres sur postgres@localhost : 5432

Éditeur SQL Constructeur graphique de requêtes

Requêtes précédentes [dropdown] [Supprimer] [Tout supprimer]

```
SELECT tablename, attname, avg_width,
       n_distinct, most_common_vals, most_common_freqs, histogram_bounds,
       correlation
FROM pg_stats;
```

Panneau sortie

Sortie de données Expliquer (Explain) Messages Historique

	QUERY PLAN text
1	Hash Join (cost=62.35..147.33 rows=5 width=466)
2	Hash Cond: ((a.attrelid = c.oid) AND (a.attnum = s.staattnum))
3	Join Filter: has_column_privilege(c.oid, a.attnum, 'select'::text)
4	-> Seq Scan on pg_attribute a (cost=0.00..67.16 rows=2316 width=70)
5	Filter: (NOT attisdropped)
6	-> Hash (cost=55.96..55.96 rows=426 width=412)
7	-> Hash Join (cost=14.84..55.96 rows=426 width=412)
8	Hash Cond: (s.starelid = c.oid)
9	-> Seq Scan on pg_statistic s (cost=0.00..35.26 rows=426 width=344)
10	-> Hash (cost=11.04..11.04 rows=304 width=72)
11	-> Seq Scan on pg_class c (cost=0.00..11.04 rows=304 width=72)

OK. Unix Ligne 4, Col 8, Caract. 141 11 lignes. 17 ms



Query – postgres sur postgres@localhost : 5432

postgres sur postgres@localhost : 5432

Éditeur SQL Constructeur graphique de requêtes

Sortie de données Expliquer (Explain) Messages Historique

QUERY PLAN
text

1	Nested Loop Left Join (cost=62.48..148.51 rows=5 width=535) (actual time=2.185..11.651 rows=426 loops=1)
2	Buffers: shared hit=935
3	-> Hash Join (cost=62.35..147.08 rows=5 width=475) (actual time=2.134..6.842 rows=426 loops=1)
4	Hash Cond: ((a.attrelid = c.oid) AND (a.attnum = s.staattnum))
5	Join Filter: has_column_privilege(c.oid, a.attnum, 'select'::text)
6	Buffers: shared hit=83
7	-> Seq Scan on pg_attribute a (cost=0.00..67.16 rows=2316 width=70) (actual time=0.012..1.897 rows=2352 loops=1)
8	Filter: (NOT attisdropped)
9	Buffers: shared hit=44
10	-> Hash (cost=55.96..55.96 rows=426 width=421) (actual time=2.077..2.077 rows=426 loops=1)
11	Buckets: 1024 Batches: 1 Memory Usage: 176kB
12	Buffers: shared hit=39
13	-> Hash Join (cost=14.84..55.96 rows=426 width=421) (actual time=0.536..1.280 rows=426 loops=1)
14	Hash Cond: (s.starelid = c.oid)
15	Buffers: shared hit=39
16	-> Seq Scan on pg_statistic s (cost=0.00..35.26 rows=426 width=349) (actual time=0.007..0.131 rows=426 loops=1)
17	Buffers: shared hit=31
18	-> Hash (cost=11.04..11.04 rows=304 width=72) (actual time=0.506..0.506 rows=304 loops=1)
19	Buckets: 1024 Batches: 1 Memory Usage: 31kB
20	Buffers: shared hit=8
21	-> Seq Scan on pg_class c (cost=0.00..11.04 rows=304 width=72) (actual time=0.006..0.204 rows=304 loops=1)
22	Buffers: shared hit=8
23	-> Index Scan using pg_namespace_oid_index on pg_namespace n (cost=0.13..0.19 rows=1 width=68) (actual time=0.004..0.005 rows=1 loops=426)
24	Index Cond: (oid = c.relnamespace)
25	Buffers: shared hit=852
26	Total runtime: 12.158 ms

OK. Unix Ligne 1, Col 1, Caract. 1 26 lignes. 28 ms

n_distinct



- If greater than zero,
 - the estimated number of distinct values in the column.
- If less than zero,
 - the negative of the number of distinct values divided by the number of rows.
 - The negated form is used when ANALYZE believes that the number of distinct values is likely to increase as the table grows;
 - The positive form is used when the column seems to have a fixed number of possible values.
 - -1 indicates a unique column in which the number of distinct values is the same as the number of rows.

Pgstats : n_distinct



```
emmanuelfuchs — psql — 79x12
cavbd=> select attname, n_distinct from pg_stats where tablename='city';
   attname   | n_distinct
-----+-----
 district    | -0.335131
 population   | -0.955381
 id           | -1
 name         | -0.980878
 countrycode  | 232
(5 rows)

cavbd=> █
```

Pgadmin III city.id statistics



The screenshot shows the pgAdmin III interface. On the left, the 'Navigateur d'objets' (Object Navigator) shows a tree structure with 'city' selected, and its columns listed: 'id', 'name', 'countrycode', 'district', and 'population'. The 'id' column is highlighted. On the right, the 'Statistiques' (Statistics) tab is active, displaying a table of statistics for the selected column.

Statistique	Valeur
Division par zéro	0
Largeur moyenne	4
Valeurs distinctes	-1
Valeurs les plus courantes	
Fréquences les plus courantes	
Limites de l'histogramme	{1,41,82,123,164,204,245,286,...
Corrélation	1

Below the statistics, the 'Panneau SQL' (SQL Panel) shows the following SQL commands:

```
-- Column: id
-- ALTER TABLE city DROP COLUMN id;
ALTER TABLE city ADD COLUMN id integer;
```

At the bottom of the window, a status bar indicates: 'Récupération des informations sur la colonne id...Exécuté. 0,08 secondes'.

Pgadmin III city.district statistics



pgAdmin III

Navigateur d'objets

- random_string(integer)
- Séquences (0)
- Tables (4)
 - city
 - Colonnes (5)
 - id
 - name
 - countrycode
 - district
 - population
 - Contraintes (1)
 - Index (0)
 - Règles (0)
 - triggers (0)
 - country
 - Colonnes (15)

Statistiques

Statistique	Valeur
Division par zéro	0
Largeur moyenne	10
Valeurs distinctes	-0.335131
Valeurs les plus courantes	{England,"São Paulo",California,Punjab,"West Bengali","Uttar Pradesh",Nordrhein-
Fréquences les plus courantes	{0.0174062,0.0169159,0.0166708,0.0115224,0.0112773,0.0105418,0.009070
Limites de l'histogramme	{"" ,Aguascalientes,Altai,Antalya,Apulia,Asir,Atyrau,Bali,Baskimaa,Bengkulu,Bolíva
Corrélation	0.00018223

Panneau SQL

```
-- Column: district  
  
-- ALTER TABLE city DROP COLUMN district;  
  
ALTER TABLE citv ADD COLUMN district text;
```

Récupération des informations sur la colonne district...Exécuté. 0,09 secondes

Pgadmin III city.name statistics



The screenshot shows the pgAdmin III interface. On the left, the 'Navigateur d'objets' (Object Navigator) pane shows a tree structure: 'random_string(integer)' > 'Séquences (0)' > 'Tables (4)' > 'city' > 'Colonnes (5)' > 'name' (selected). Below this, other objects like 'countrycode', 'district', 'population', 'Contraintes (1)', 'Index (0)', 'Règles (0)', 'triggers (0)', and 'country' are visible.

The main pane is titled 'Statistiques' (Statistics) and displays a table of statistics for the selected column 'name'.

Statistique	Valeur
Division par zéro	0
Largeur moyenne	9
Valeurs distinctes	-0.980878
Valeurs les plus courantes	{Cambridge,Córdoba,Hamilton,"La Paz",León,Richmond,"San Fernando","San Jose
Fréquences les plus courantes	{0.000735474,0.000735474,0.000735474,0.000735474,0.000735474,0.00073
Limites de l'histogramme	{"A Coruña (La Coruña)",Agra,Almaty,"Andorra la Vella",Arak,Athenai,Bairiki,Baqu
Corrélation	0.0292628

Below the statistics table is the 'Panneau SQL' (SQL Panel) containing the following SQL code:

```
-- Column: name  
  
-- ALTER TABLE city DROP COLUMN name;  
  
ALTER TABLE citv ADD COLUMN name text;
```

The status bar at the bottom indicates 'Récupération des informations sur la colonne name...Exécuté.' and '0,14 secondes'.

Pgadmin III city.countrycode statistics



pgAdmin III

Navigateur d'objets

- random_string(integer)
- Séquences (0)
- Tables (4)
 - city
 - Colonnes (5)
 - id
 - name
 - countrycode
 - district
 - population
 - Contraintes (1)
 - Index (0)
 - Règles (0)
 - triggers (0)
 - country
 - Colonnes (15)

Statistiques

Statistique	Valeur
Division par zéro	0
Largeur moyenne	4
Valeurs distinctes	232
Valeurs les plus courantes	{CHN,IND,USA,BRA,JPN,RUS,MEX,PHL,DEU,IDN,GBR,KOR,IRN,NGA,TUR,ESP,PAK,ITA}
Fréquences les plus courantes	{0.0889924,0.0835989,0.0671733,0.0612895,0.0607992,0.0463349,0.042412}
Limites de l'histogramme	{ABW,AGO,ARE,ARM,AUS,AUS,AUT,BEL,BEN,BGR,BGR,BLR,BLR,BLZ,BOL,BWA,CHE,C}
Corrélation	0.668009

Panneau SQL

```
-- Column: countrycode  
  
-- ALTER TABLE city DROP COLUMN countrycode;  
  
ALTER TABLE city ADD COLUMN countrycode character(3);
```

n_distinct r = 11



```
emmanuelfuchs — psql — 75x7
cavbd=> select attname, n_distinct from pg_stats where tablename='r';
 attname | n_distinct 
-----+-----
 ri      |          11
(1 row)

cavbd=>
```

n_distinct r = 11



```
emmanuel fuchs — psql — 126x21
cavbd=> \x
Expanded display is on.
cavbd=> select * from pg_stats where tablename='r';
-[ RECORD 1 ]-----
 schemaname | public
 tablename  | r
 attname    | ri
 inherited  | f
 null_frac  | 0
 avg_width  | 4
 n_distinct | 11
 most_common_vals | {6,7,3,8,4,5,2,9,1,10,0}
 most_common_freqs | {0.102433,0.1017,0.100867,0.1008,0.100433,0.0987667,0.0983667,0.0979,0.0974667,0.0522333,0.0490333}
 histogram_bounds | 
 correlation  | 0.0961845
 most_common_elems | 
 most_common_elem_freqs | 
 elem_count_histogram | 
cavbd=> █
```

most_common_freqs



- A list of the frequencies of the most common values, i.e., number of occurrences of each divided by total number of rows.

most_common_value



```
emmanuel fuchs — psql — 126x21
cavbd=> \x
Expanded display is on.
cavbd=> select * from pg_stats where tablename='r';
-[ RECORD 1 ]-----
 schemaname | public
 tablename  | r
 attname    | ri
 inherited  | f
 null_frac  | 0
 avg_width  | 4
 n_distinct | 11

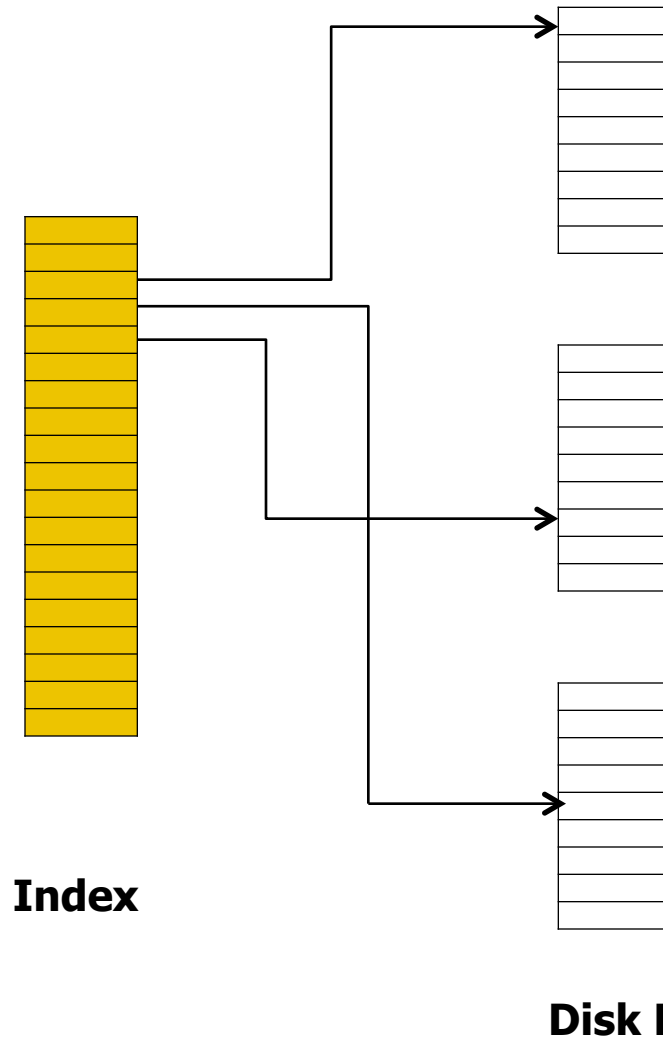
n_distinct | 11
most_common_vals | {6,7,3,8,4,5,2,9,1,10,0}
most_common_freqs | {0.102433,0.1017,0.100867,0.100867,0.100867,0.100867,0.100867,0.100867,0.100867,0.100867,0.100867}
```

correlation




- Statistical correlation between physical row ordering and logical ordering of the column values.
- This ranges from -1 to +1.
- When the value is near -1 or +1, an index scan on the column will be estimated to be cheaper than when it is near zero, due to reduction of random access to the disk.
- This column is NULL if the column data type does not have a < operator.

Index and disk blocks correlation



0

Index and disk blocks correlation



1

