


Conception Avancée de Bases de Données

ANALYZE



ANALYZE



- ANALYZE collects statistics about the contents of tables in the database, and stores the results in the pg_statistic system catalog.
- The query planner uses these statistics to help determine the most efficient execution plans for queries.



ANALYZE Behavior



- Statistics data are not updated on the fly.
- ANALYSE (as well VACUUM) command forces statistics update.
- With no parameter, ANALYZE examines every table in the current database.
- With a parameter, ANALYZE examines only that table.

R creation



```
emmanuelfuchs — psql — 63x5
cavbd=> CREATE TABLE r (ri integer);
CREATE TABLE
cavbd=> █
```

```
emmanuelfuchs — psql — 87x5
cavbd=> CREATE TABLE r (ri integer);
CREATE TABLE
cavbd=> INSERT INTO r (ri) SELECT round(random()*10) FROM generate_series(1, 1000000);
INSERT 0 1000000
cavbd=> █
```

pg_class not updated



```
emmanuelfuchs — psql — 97x12
cavbd=> CREATE TABLE r (ri integer);
CREATE TABLE
cavbd=> INSERT INTO r (ri) SELECT round(random()*10) FROM generate_series(1, 1000000);
INSERT 0 1000000
cavbd=> SELECT relname, reltuples, relpages FROM pg_catalog.pg_class WHERE relname='r';
 relname | reltuples | relpages 
-----+-----+-----
 r       |         0 |         0
(1 row)

cavbd=> █
```

pg_class location



Query – cavbd sur postgres@localhost : 5432 *

Éditeur SQL Constructeur graphique de requêtes

Requêtes précédentes Supprimer Tout supprimer

```
SELECT relname
FROM pg_class
where relname = 'r';
```

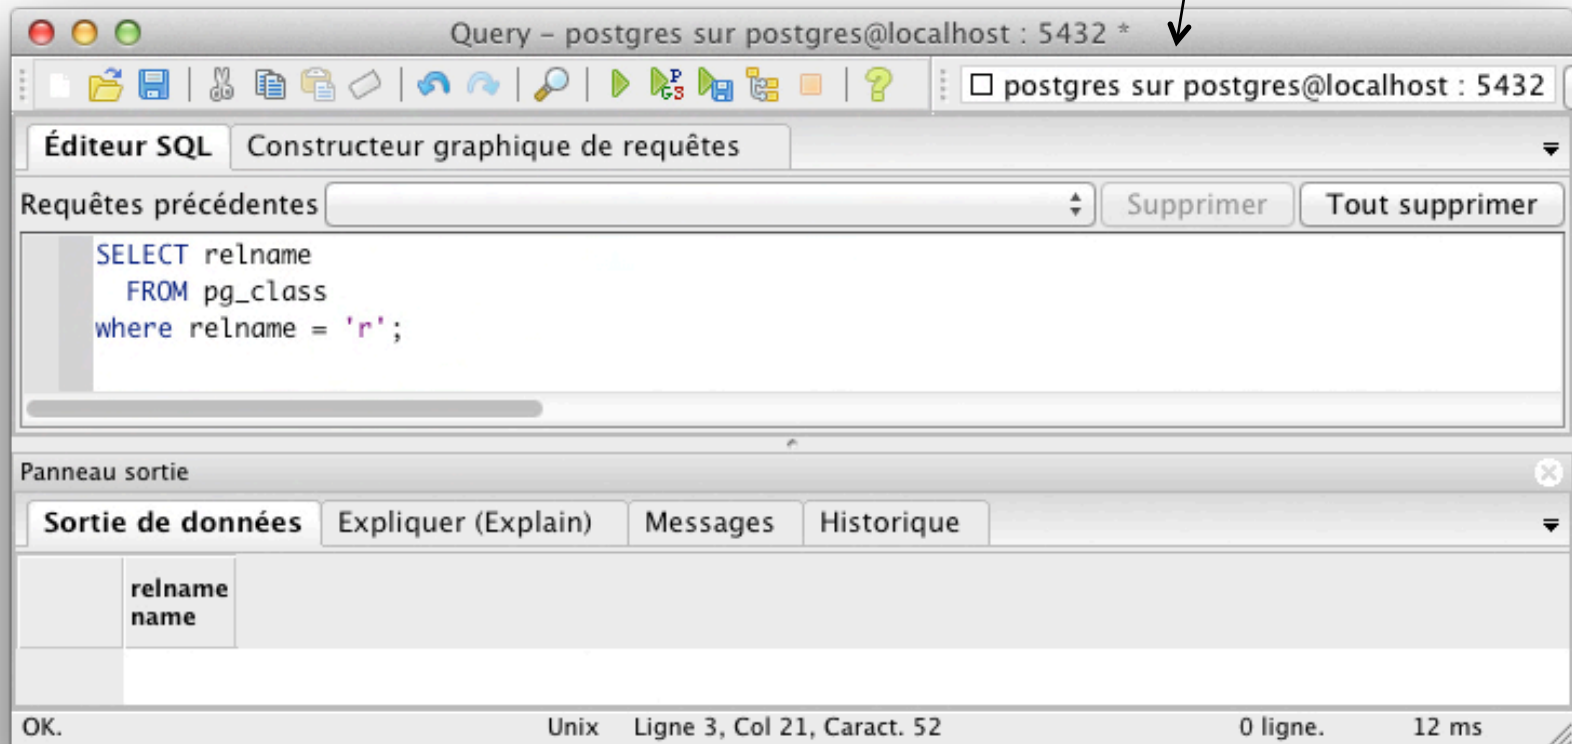
Panneau sortie

Sortie de données Expliquer (Explain) Messages Historique

	relname
1	r

OK. Unix Ligne 3, Col 21, Caract. 52 1 ligne. 12 ms

Attention au catalog



pg_stats view not updated



```
emmanuelfuchs — psql — 71x25
cavbd=> CREATE TABLE r (ri integer);
CREATE TABLE
cavbd=> INSERT INTO r (ri) SELECT round(random()*10) FROM generate_series(1, 1000000);
INSERT 0 1000000
cavbd=> SELECT relname, reltuples, relpages FROM pg_catalog.pg_class WHERE relname='r';
 relname | reltuples | relpages 
-----+-----+-----
 r       |          0 |          0
(1 row)

cavbd=> select * from cavbd.pg_catalog.pg_stats where tablename='r';
 schemaname | tablename | attname | inherited | null_frac | avg_width |
n_distinct | most_common_vals | most_common_freqs | histogram_bounds |
correlation | most_common_elems | most_common_elem_freqs | elem_count_
histogram
-----+-----+-----+-----+-----+-----+
(0 rows)

cavbd=> █
```


pg_stats view not updated



```
emmanuelfuchs — psql — 71x25
cavbd=> CREATE TABLE r (ri integer);
CREATE TABLE
cavbd=> INSERT INTO r (ri) SELECT round(random()*10) FROM generate_series(1, 1000000);
INSERT 0 1000000
cavbd=> SELECT relname, reltuples, relpages FROM pg_catalog.pg_class WHERE relname='r';
 relname | reltuples | relpages
-----+-----+-----
 r       |          0 |          0
```

```
(1 row)
```

reltuples and relpages are not updated

```
cavbd=> select * from cavbd.pg_catalog.pg_stats where tablename='r';
 schemaname | tablename | attname | inherited | null_frac | avg_width |
n_distinct | most_common_vals | most_common_freqs | histogram_bounds |
correlation | most_common_elems | most_common_elem_freqs | elem_count_
histogram
```

```
(0 rows)
```

```
cavbd=>
```

After ANALYZE pg_class is updated



```
emmanuelfuchs — psql — 93x27
cavbd=> CREATE TABLE r (ri integer);
CREATE TABLE
cavbd=> INSERT INTO r (ri) SELECT round(random()*10) FROM generate_series(1, 1000000);
INSERT 0 1000000
cavbd=> SELECT relname, reltuples, relpages FROM pg_catalog.pg_class WHERE relname='r';
 relname | reltuples | relpages 
-----+-----+-----
 r      |          0 |          0
(1 row)

cavbd=> select * from cavbd.pg_catalog.pg_stats where tablename='r';
 schemaname | tablename | attname | inherited | null_frac | avg_width | n_distinct | most_com
mon_vals | most_common_freqs | histogram_bounds | correlation | most_common_elems | most_comm
on_elem_freqs | elem_count_histogram
-----+-----+-----+-----+-----+-----+-----+-----
(0 rows)

cavbd=> ANALYZE r;
ANALYZE
cavbd=> SELECT relname, reltuples, relpages FROM pg_catalog.pg_class WHERE relname='r';
 relname | reltuples | relpages 
-----+-----+-----
 r      |      1e+06 |      4425
(1 row)

cavbd=>
```

After ANALYZE pg_stats view is updated



```
emmanuelfuchs — psql — 91x31
cavbd=> CREATE TABLE r (ri integer);
CREATE TABLE
cavbd=> INSERT INTO r (ri) SELECT round(random()*10) FROM generate_series(1, 1000000);
INSERT 0 1000000
cavbd=> select attname from cavbd.pg_catalog.pg_stats where tablename='r';
 attname
-----
(0 rows)

cavbd=> SELECT relname, reltuples, relpages FROM pg_catalog.pg_class WHERE relname='r';
 relname | reltuples | relpages
-----+-----+-----
 r       |          0 |          0
(1 row)

cavbd=> ANALYZE r;
ANALYZE
cavbd=> SELECT relname, reltuples, relpages FROM pg_catalog.pg_class WHERE relname='r';
 relname | reltuples | relpages
-----+-----+-----
 r       |      1e+06 |        4425
(1 row)

cavbd=> select attname from cavbd.pg_catalog.pg_stats where tablename='r';
 attname
-----
 ri
(1 row)

cavbd=> █
```

After ANALYZE pg_stats view is updated



```
emmanuelfuchs — psql — 98x34
cavbd=> CREATE TABLE r (ri integer);
CREATE TABLE
cavbd=> INSERT INTO r (ri) SELECT round(random()*10) FROM generate_series(1, 1000000);
INSERT 0 1000000
cavbd=> SELECT relname, reltuples, relpages FROM pg_catalog.pg_class WHERE relname='r';
 relname | reltuples | relpages 
-----+-----+-----
 r       |         0 |         0
(1 row)

cavbd=> select * from cavbd.pg_catalog.pg_stats where tablename='r';
 schemaname | tablename | attname | inherited | null_frac | avg_width | n_distinct | most_common_v
als | most_common_freqs | histogram_bounds | correlation | most_common_elems | most_common_elem_fr
eqs | elem_count_histogram
-----+-----+-----+-----+-----+-----+-----+-----
(0 rows)

cavbd=> ANALYZE r;
ANALYZE
cavbd=> SELECT relname, reltuples, relpages FROM pg_catalog.pg_class WHERE relname='r';
 relname | reltuples | relpages 
-----+-----+-----
 r       |      1e+06 |      4425
(1 row)

cavbd=> select attname, n_distinct from cavbd.pg_catalog.pg_stats where tablename='r';
 attname | n_distinct 
-----+-----
 ri      |         11
(1 row)

cavbd=>
```

