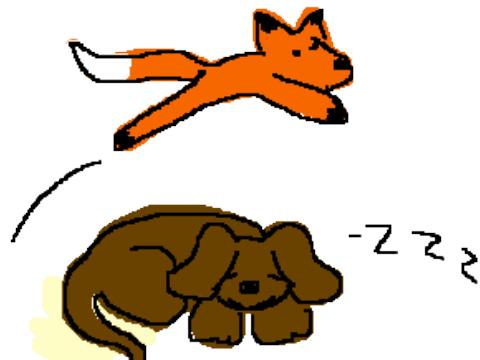
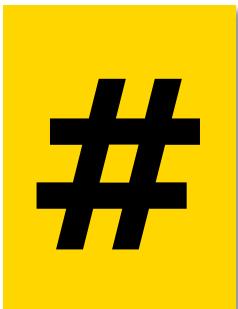
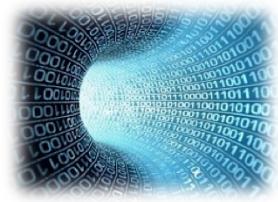


MAP REDUCE

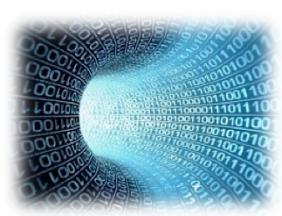


Big Data Definition



- Big Data is a collection of large datasets that cannot be processed using traditional computing techniques.
 - Any data that can challenge our current technology

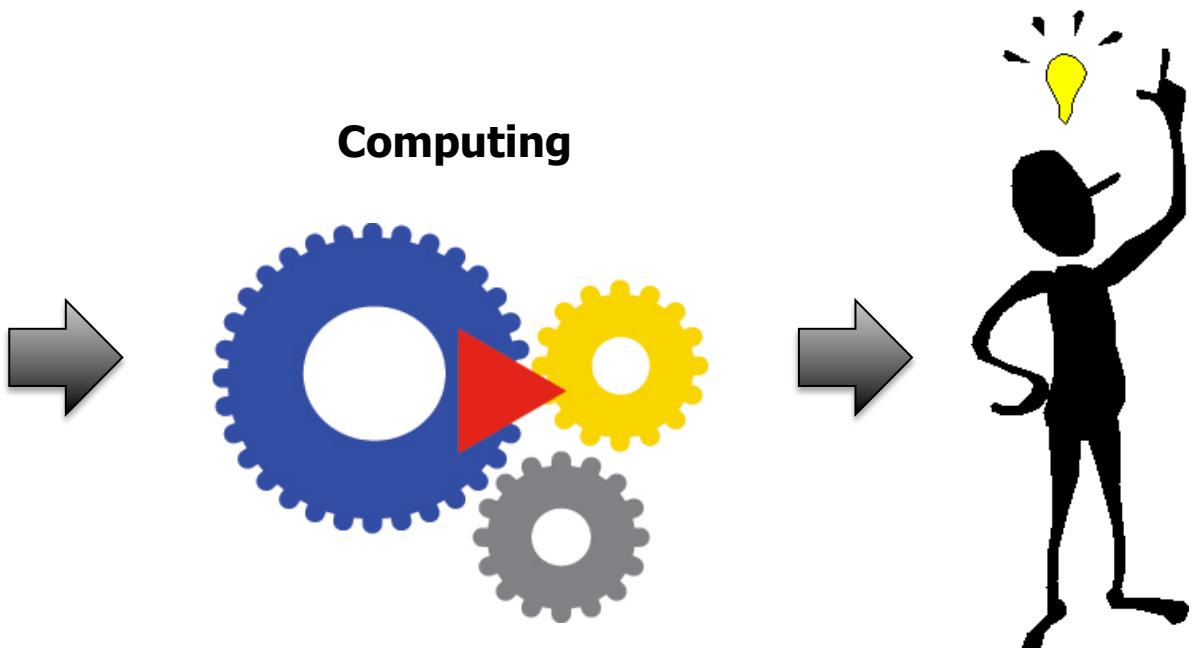
Big Data

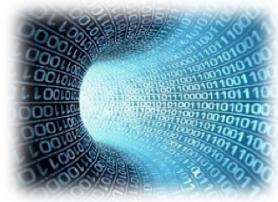


Big data is about the Insight that we want to extract from information.

Data

Computing

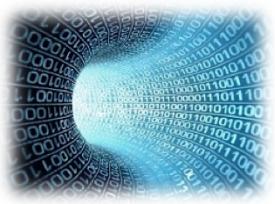




Big Data Frameworks properties

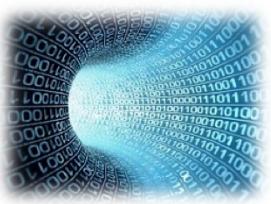
- Big Data Frameworks properties
 - Elasticity
 - Errors Transparency
 - Scalability





Abstraction

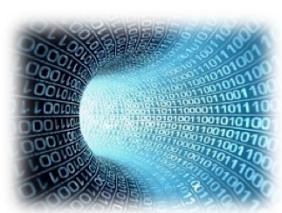
- The primary MapReduce abstraction is parallelizable, easy to understand, and hides the details of distributed computing, thus allowing Big Data engine to guarantee correctness.
- Specification of computation as a mapping then a reduction.



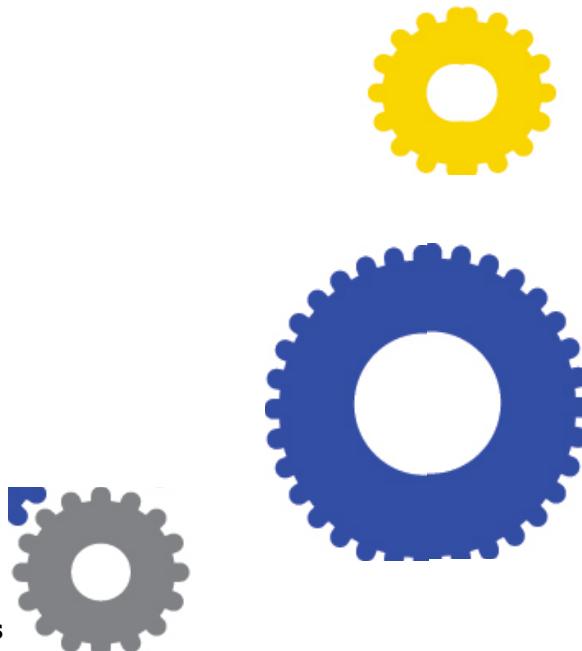
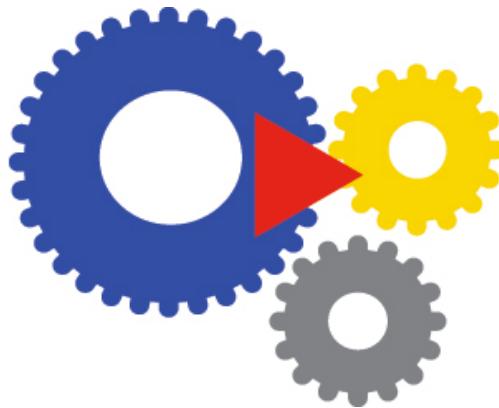
MapReduce

- Map Reduce
 - Inspired from functional programming paradigms.
- Two functions
 - Map : specify the per-record computation
 - Reduce : specify result aggregation
- MapReduce :
 - Theoretical framework.
- Hadoop MapReduce Framework :
 - open source implementation

MapReduce



- Re-think the best approach to an old problem.
- Need to “unlearn” things in order to full utilize the power of a framework.
- There may be even better solutions that we didn’t discovered yet.



Interface & Encapsulation = New paradigm



Carroserie Hypomobile



Domain Maturity



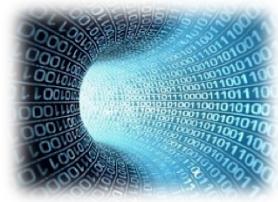
Traduction en cours

Back to the past !!!

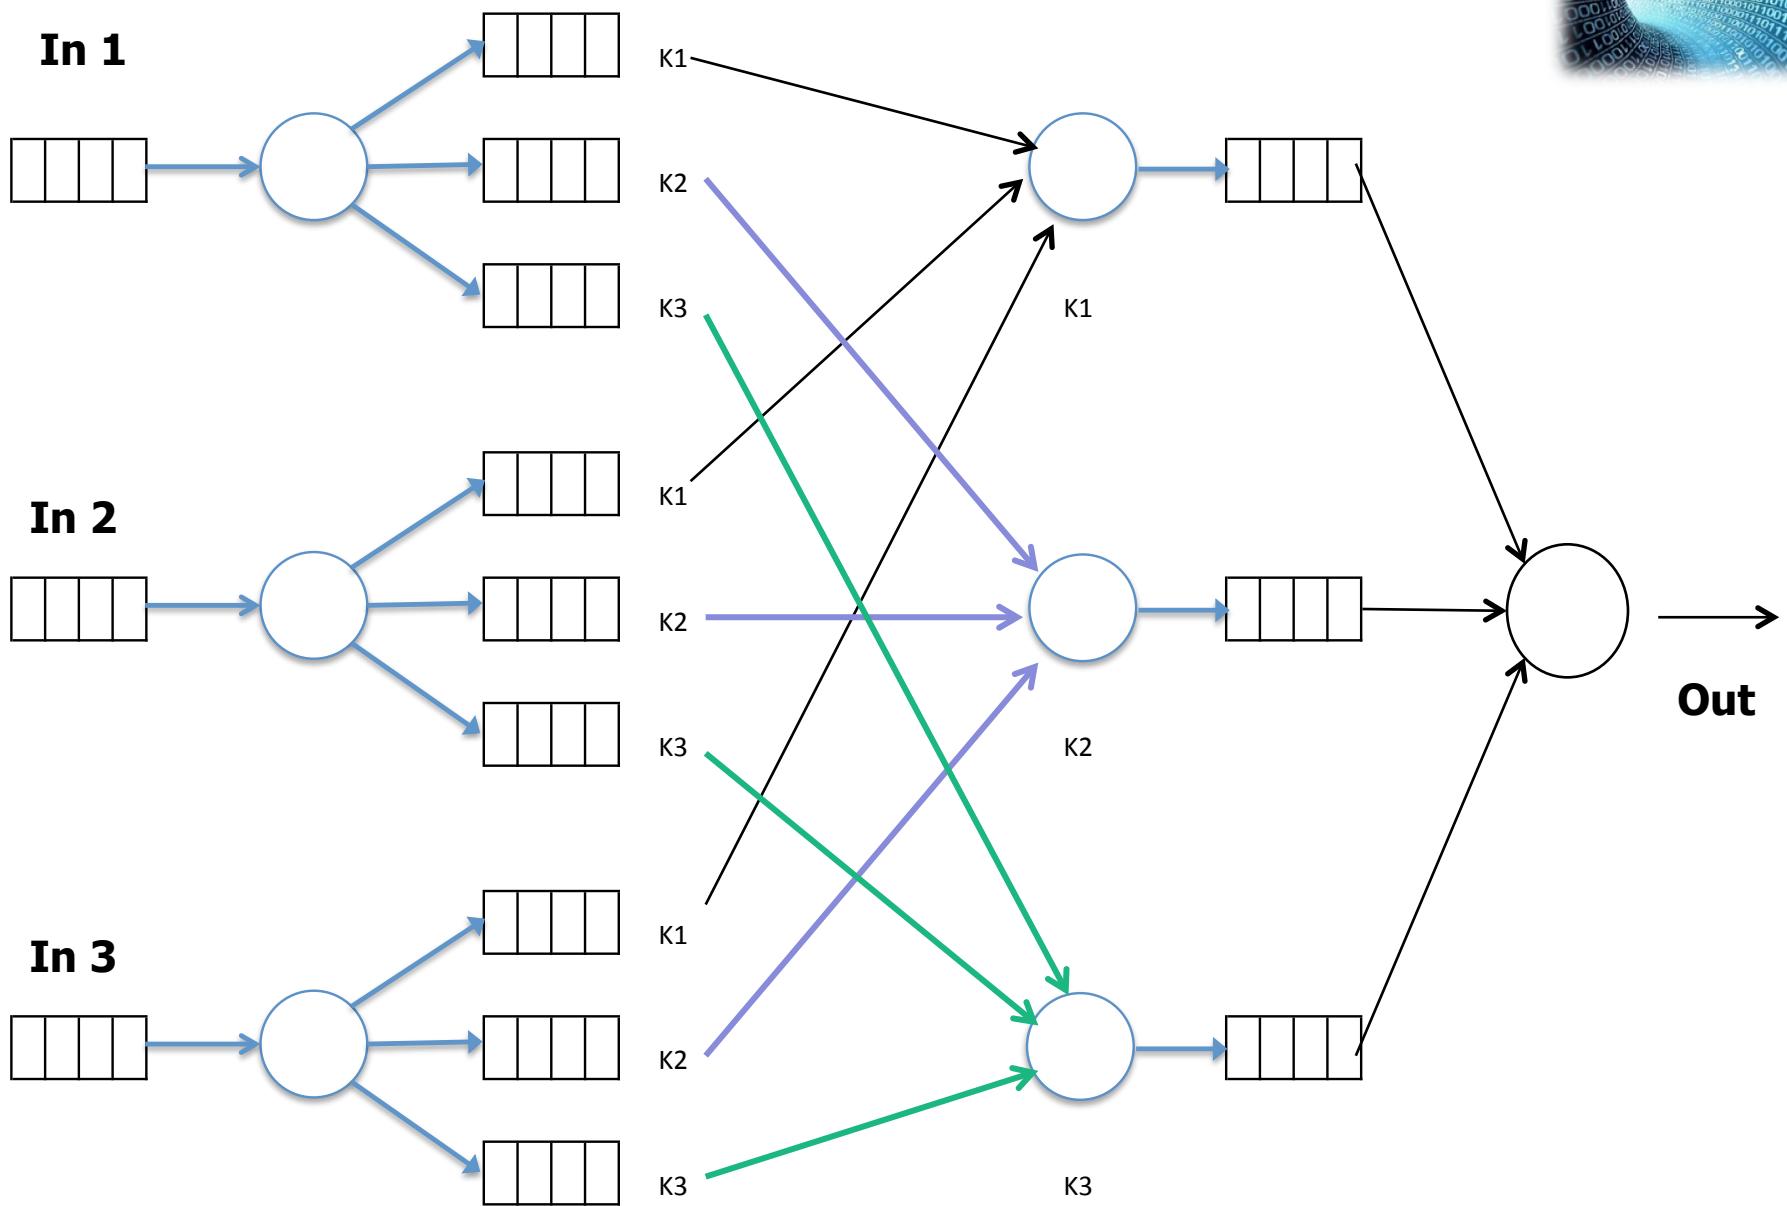
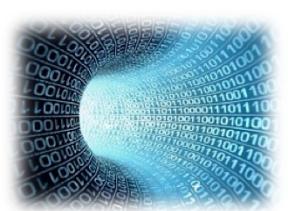


MAP REDUCE

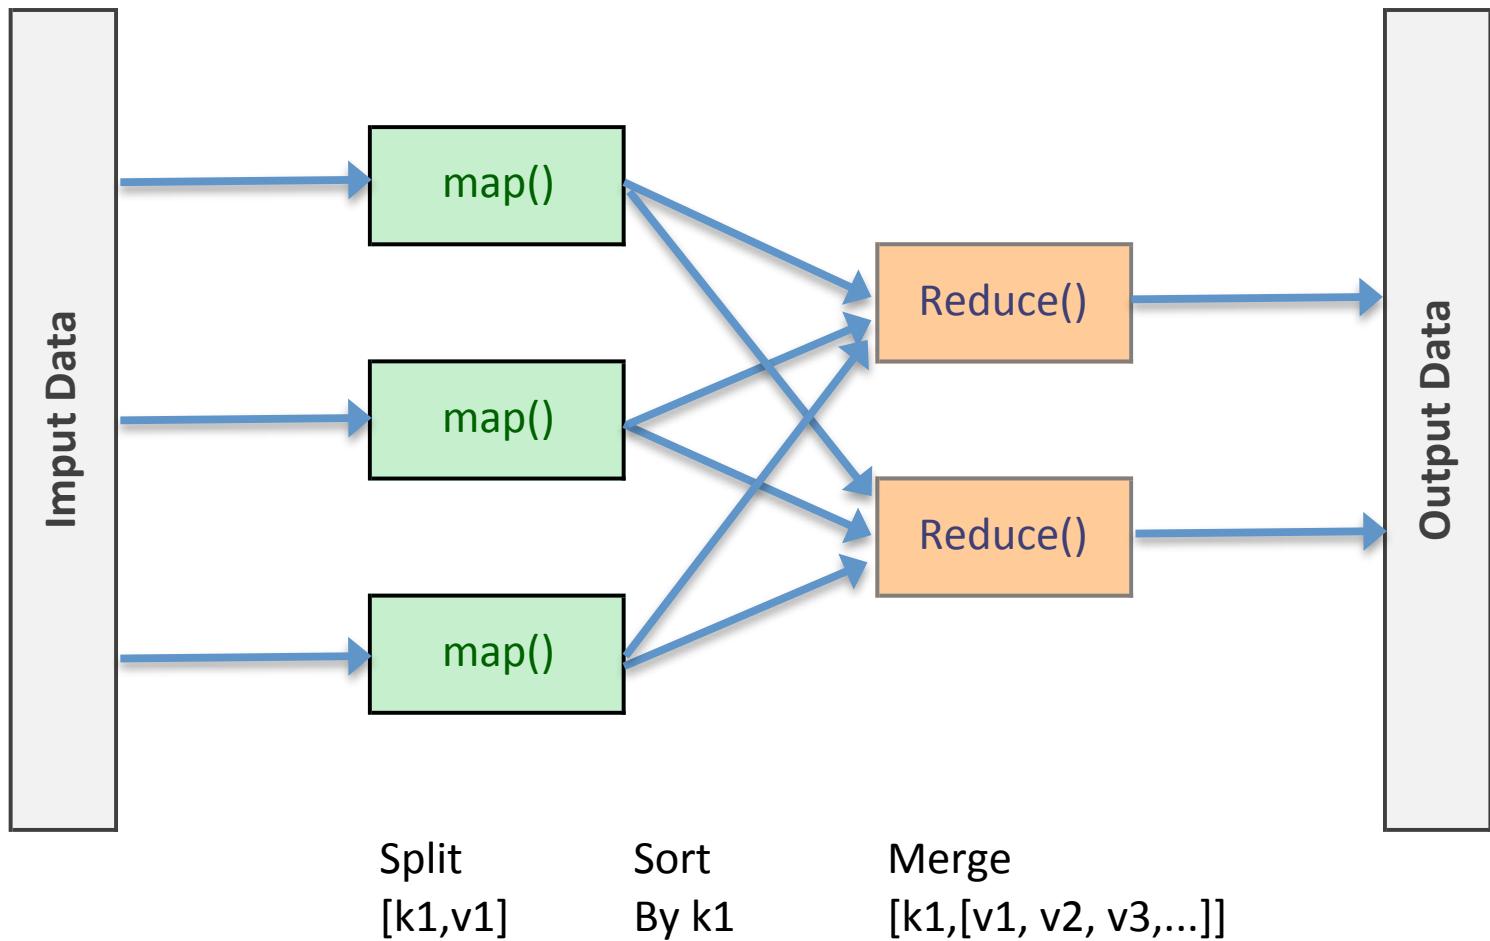
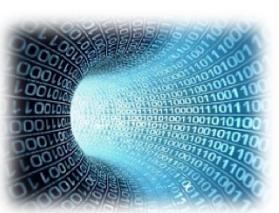
<http://martinfowler.com/articles/nosqlKeyPoints.html>

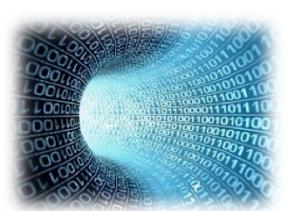


- The map task reads data from an aggregate and boils it down to relevant key-value pairs.
- Maps only read a single record at a time and can thus be parallelized and run on the node that stores the record.
- Reduce tasks take many values for a single key output from map tasks and summarize them into a single output.
- Each reducer operates on the result of a single key, so it can be parallelized by key.



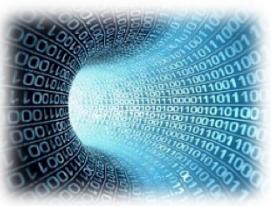
MapReduce





if you have a hammer, everything looks like a nail.



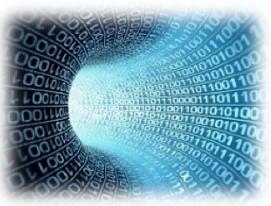


Map Reduce

- Transform lists of input data elements into lists of output data elements.
- Twice with two different list processsing approach :
 - 1) MAP : Partitioning
 - 2) Reduce : Merge Select (join)

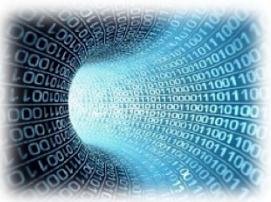


Traduction en cours



Lists

- Key Value Lists.
 - 1) MAP : Partitioning on Key
 - 2) Reduce : Merge operation on Value

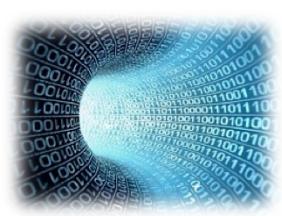


Map Reduce Canonical Exemple

- How many times different words appear in a set of sentences ?

Counting Words

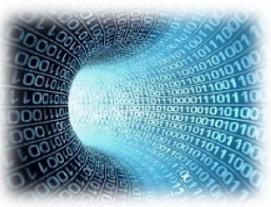




Pangrams

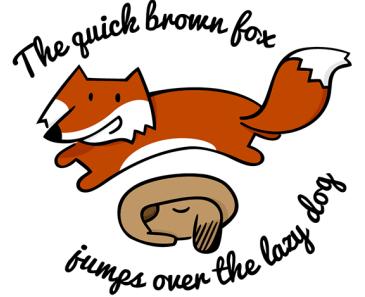
- Pangrams :
 - The quick brown fox jumps over the lazy dog.
 - (35 letters)
- A pangram, or holoalphabetic sentence, includes every letter of the alphabet at least once.
- "Portez ce vieux whisky au juge blond qui fume"





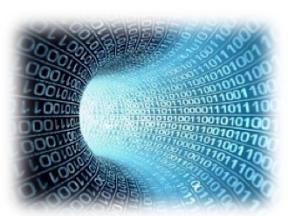
The Fox

The
quick
brown
fox
jumps
over
the
lazy
dog.



Map

Key	Value
The	1
quick	1
brown	1
fox	1
jumps	1
over	1
the	1
lazy	1
dog.	1

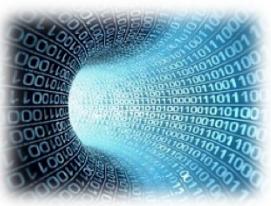


Map

Key	Value
The	1
quick	1
brown	1
fox	1
jumps	1
over	1
the	1
lazy	1
dog.	1

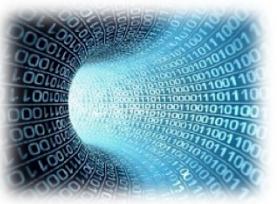


- The map task reads data from an aggregate and boils it down to relevant key-value pairs.



Sentences

- Dogs That Look Like a Fox
- How to Stop Your Dog Jumping on Guests
- Dog Who Looks Like A Blue Eyed Fox



Map

Dogs
That
Look
Like
a
Fox

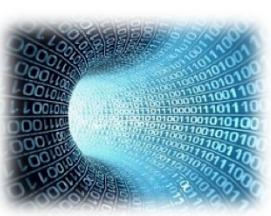
How
to
Stop
Your
Dog
Jumping
on
Guests

Dog
Who
Looks
Like
A
Blue
Eyed
Fox

Key	Value
Dogs	1
That	1
Look	1
Like	1
a	1
Fox	1

Key	Value
How	1
to	1
Stop	1
Your	1
Dog	1
Jumping	1
on	1
Guests	1

Key	Value
Dog	1
Who	1
Looks	1
Like	1
A	1
Blue	1
Eyed	1
Fox	1

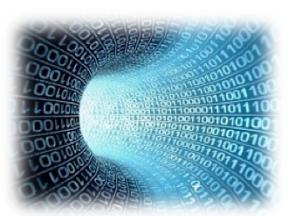


Map 1

Key	Value
brown	1
dog.	1
fox	1
jumps	1
lazy	1
over	1
quick	1
The	1
the	1

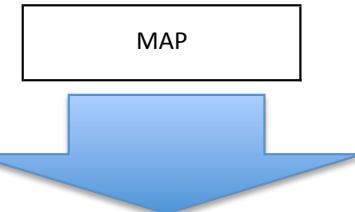


brown	dog.	fox	jumps	lazy	over	quick	The	the
1	1	1	1	1	1	1	1	1



Map 1

Key	Value
brown	1
dog.	1
fox	1
jumps	1
lazy	1
over	1
quick	1
The	1
the	1



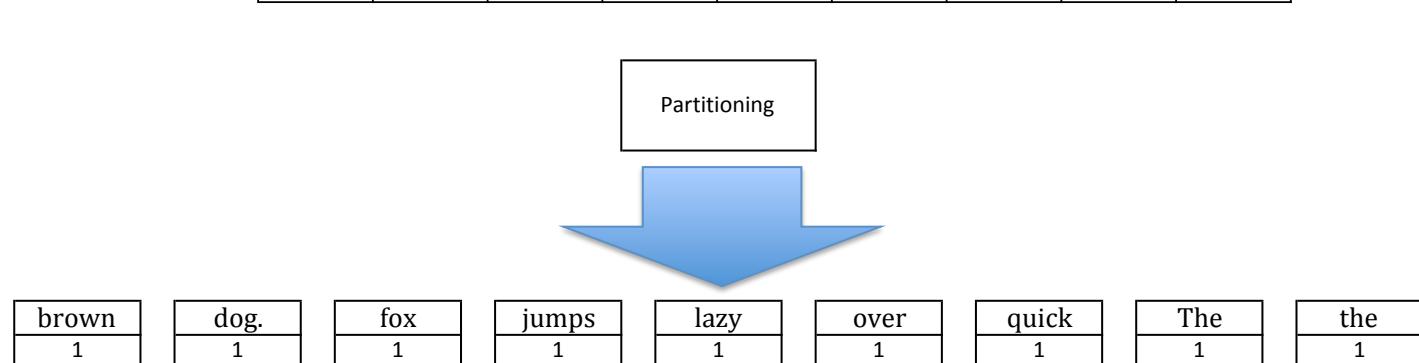
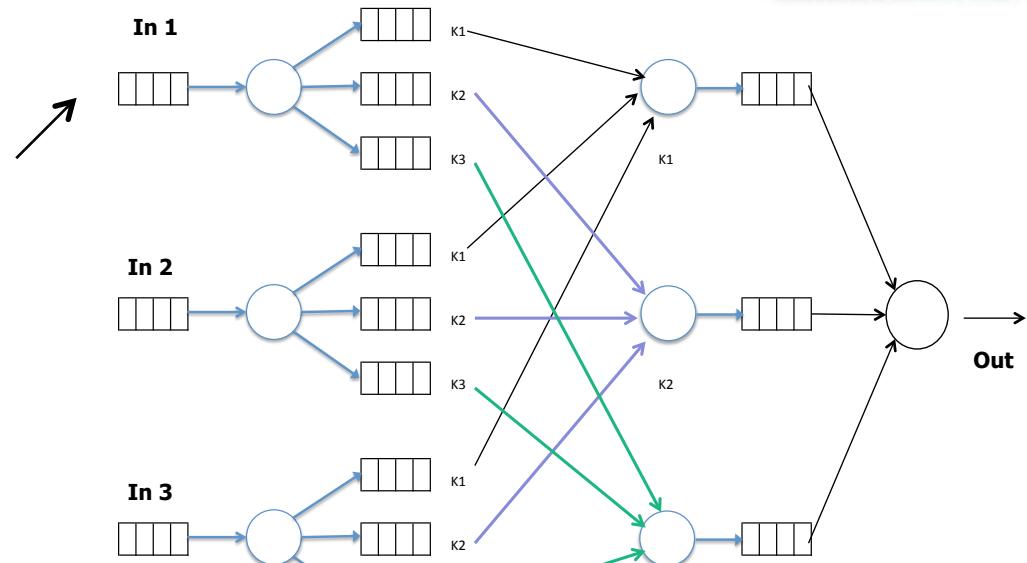
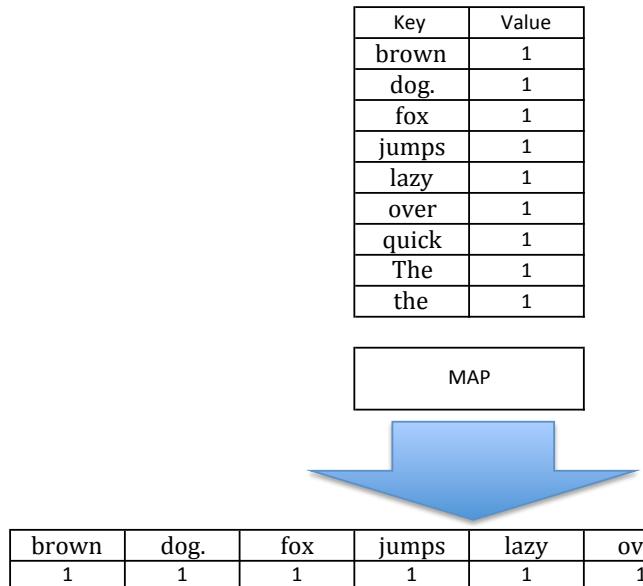
brown	dog.	fox	jumps	lazy	over	quick	The	the
1	1	1	1	1	1	1	1	1



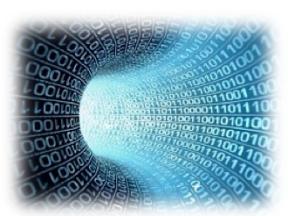
brown	dog.	fox	jumps	lazy	over	quick	The	the
1	1	1	1	1	1	1	1	1



Map 1



Map 2



Key	Value
That	1
Look	1
Like	1
Fox	1
Dogs	1
a	1

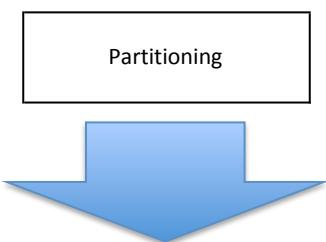
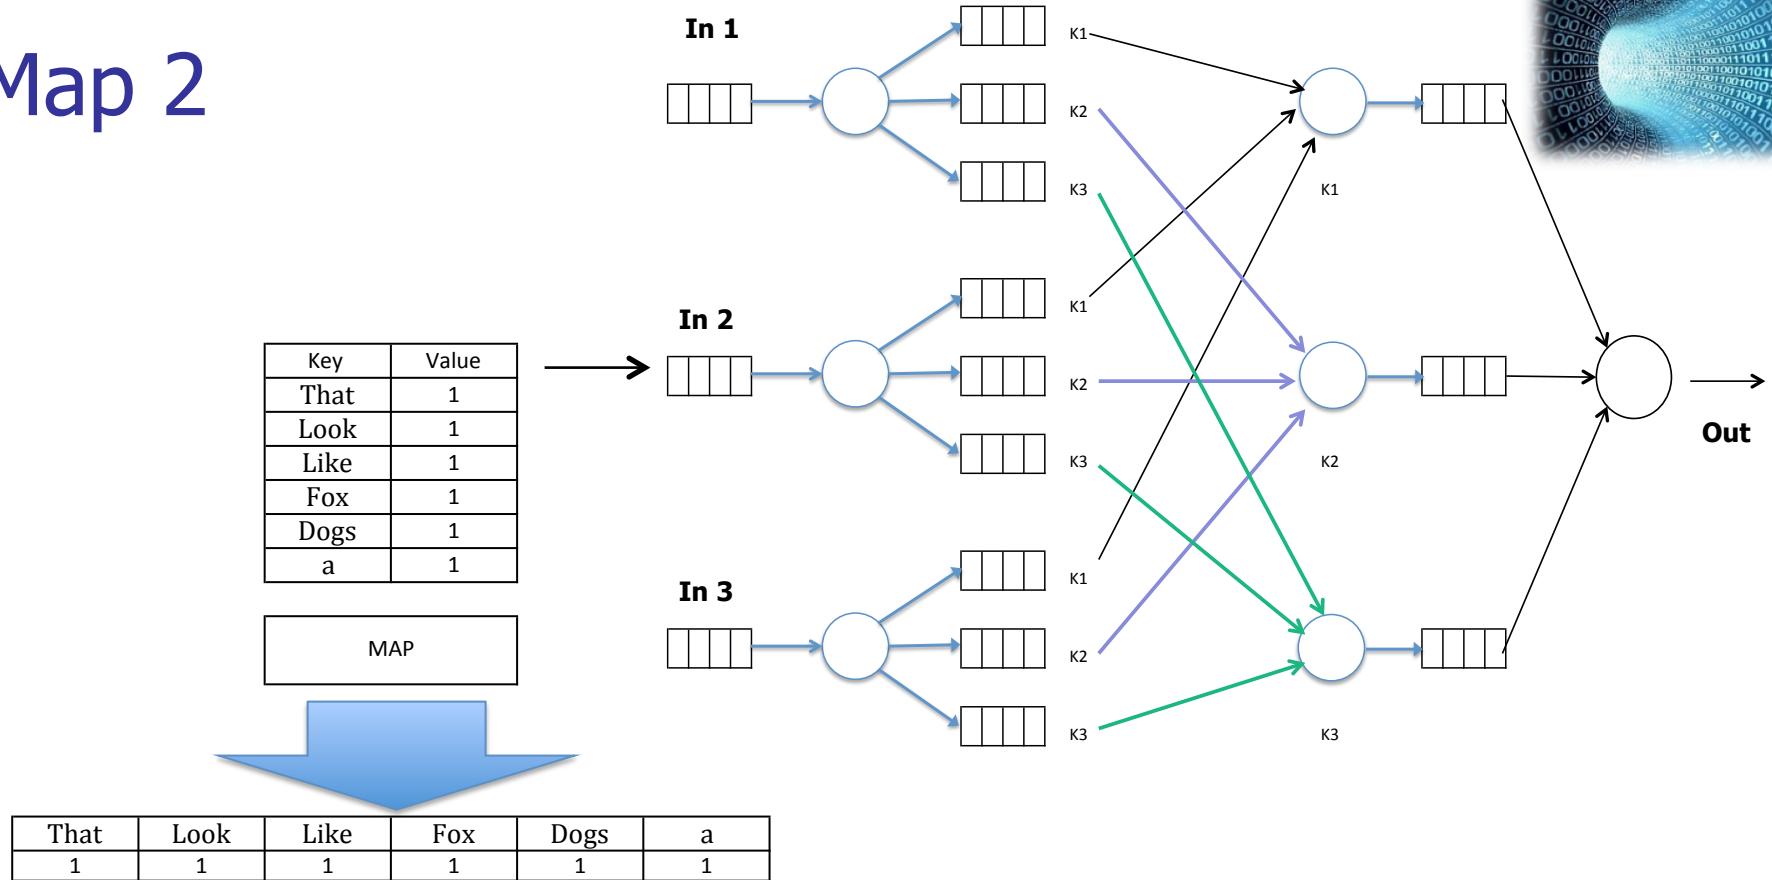
MAP

That	Look	Like	Fox	Dogs	a
1	1	1	1	1	1

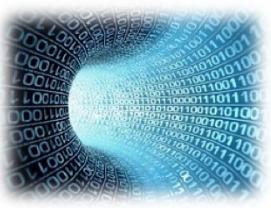
Partitioning

That	Look	Like	Fox	Dogs	a
1	1	1	1	1	1

Map 2



That	Look	Like	Fox	Dogs	a
1	1	1	1	1	1



Map 3

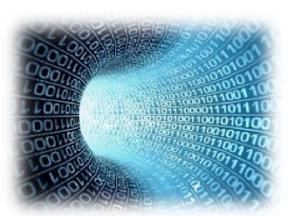
Key	Value
Dog	1
Guests	1
How	1
Jumping	1
on	1
Stop	1
to	1
Your	1

MAP

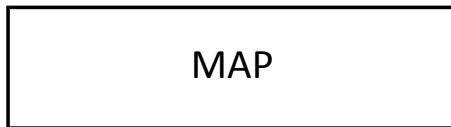


Dog	Guests	How	Jumping	on	Stop	to	Your
1	1	1	1	1	1	1	1

Map 4

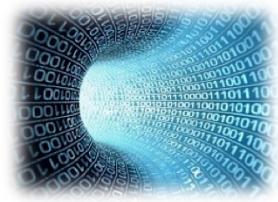


Key	Value
Who	1
Looks	1
Like	1
Fox	1
Dog	1
Blue	1
A	1
Eyed	1



Who	Looks	Like	Fox	Dog	Blue	A	Eyed
1	1	1	1	1	1	1	1

Map



Key	Value
brown	1
dog	1
fox	1
jumps	1
lazy	1
over	1
quick	1
The	1
the	1

MAP



brown	dog	fox	jumps	lazy	over	quick	The	the
1	1	1	1	1	1	1	1	1

Key	Value
That	1
Look	1
Like	1
Fox	1
Dogs	1
a	1

MAP



That	Look	Like	Fox	Dogs	a
1	1	1	1	1	1

Key	Value
Dog	1
Guests	1
How	1
Jumping	1
on	1
Stop	1
to	1
Your	1

MAP



Dog	Guests	How	Jumping	on	Stop	to	Your
1	1	1	1	1	1	1	1

Key	Value
Who	1
Looks	1
Like	1
Fox	1
Dog	1
Blue	1
A	1
Eyed	1

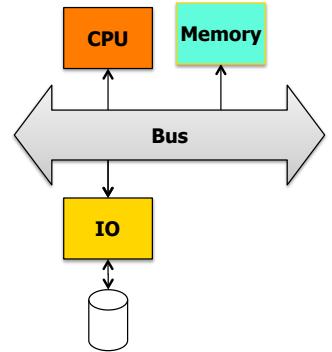
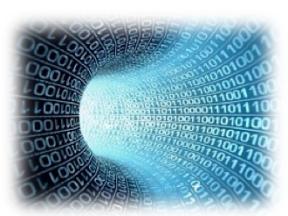
MAP



Who	Looks	Like	Fox	Dog	Blue	A	Eyed
1	1	1	1	1	1	1	1

- Maps only read a single record at a time and can thus be parallelized and run on the node that stores the record.

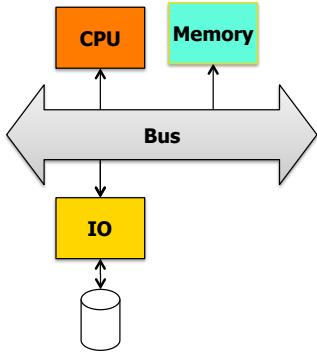
Map



Key	Value
brown	1
dog	1
fox	1
jumps	1
lazy	1
over	1
quick	1
The	1
the	1

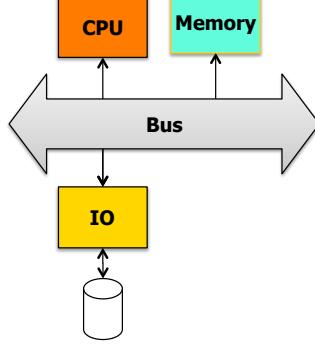
MAP

brown	dog	fox	jumps	lazy	over	quick	The	the
1	1	1	1	1	1	1	1	1



Key	Value
That	1
Look	1
Like	1
Fox	1
Dogs	1
a	1

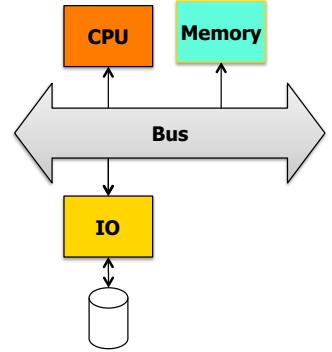
That	Look	Like	Fox	Dogs	a
1	1	1	1	1	1



Key	Value
Dog	1
Guests	1
How	1
Jumping	1
on	1
Stop	1
to	1
Your	1

MAP

Dog	Guests	How	Jumping	on	Stop	to	Your
1	1	1	1	1	1	1	1



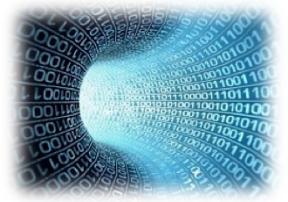
Key	Value
Who	1
Looks	1
Like	1
Fox	1
Dog	1
Blue	1
A	1
Eyed	1

MAP

Who	Looks	Like	Fox	Dog	Blue	A	Eyed
1	1	1	1	1	1	1	1

- Maps only read a single record at a time and can thus be parallelized and run on the node that stores the record.

Map



Map

Key	Value
brown	1
dog.	1
fox	1
jumps	1
lazy	1
over	1
quick	1
The	1
the	1



brown	dog	fox	jumps	lazy	over	quick	The	the
1	1	1	1	1	1	1	1	1

Key	Value
That	1
Look	1
Like	1
Fox	1
Dogs	1
a	1

MAP

That	Look	Like	Fox	Dogs	a
1	1	1	1	1	1

Key	Value
Dog	1
Guests	1
How	1
Jumping	1
on	1
Stop	1
to	1
Your	1

MAP

Dog	Guests	How	Jumping	on	Stop	to	Your
1	1	1	1	1	1	1	1

Key	Value
Who	1
Looks	1
Like	1
Fox	1
Dog	1
Blue	1
A	1
Eyed	1

MAP

Who	Looks	Like	Fox	Dog	Blue	A	Eyed
1	1	1	1	1	1	1	1

Partition

Partitioning

brown	dog	fox	jumps	lazy	over	quick	The	the
1	1	1	1	1	1	1	1	1

Partitioning

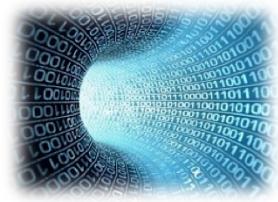
That	Look	Like	Fox	Dogs	a
1	1	1	1	1	1

Partitioning

Dog	Guests	How	Jumping	on	Stop	to	Your
1	1	1	1	1	1	1	1

Partitioning

Who	Looks	Like	Fox	Dog	Blue	A	Eyed
1	1	1	1	1	1	1	1



Key Partitioning

Fox	1
Fox	1
Fox	1

A	1
A	1

Like	1
Like	1

Your	1
------	---

Who	1
-----	---

to	1
----	---

The	1
-----	---

the	1
-----	---

That	1
------	---

Stop	1
------	---

quick	1
-------	---

over	1
------	---

on	1
----	---

Looks	1
-------	---

Look	1
------	---

lazy	1
------	---

jumps	1
-------	---

Jumping	1
---------	---

How	1
-----	---

Guests	1
--------	---

Dogs	1
------	---

brown	1
-------	---

Blue	1
------	---

Eyed	1
------	---

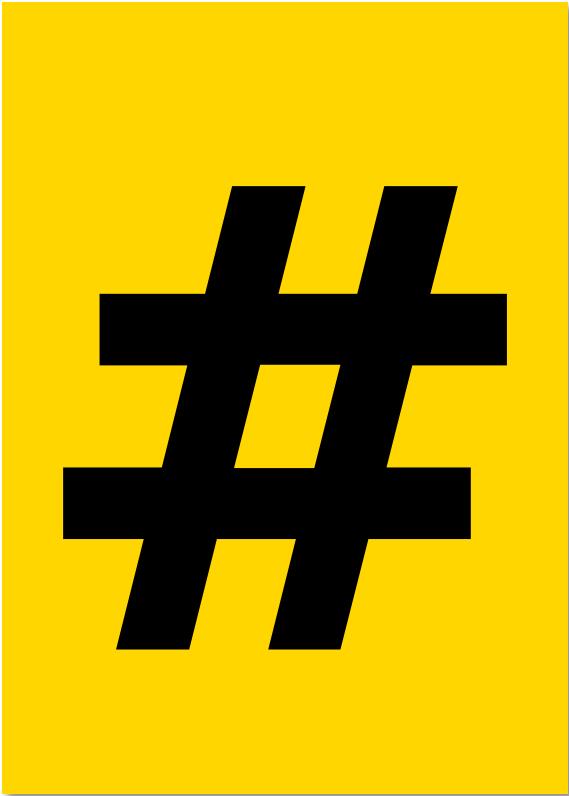
Dog	1
Dog	1
Dog	1

FOX	1
FOX	1
FOX	1

A	1
A	1

Like	1
Like	1

Key Partitioning



Fox	1
Fox	1
Fox	1

A	1
A	1

Like	1
Like	1

Your	1
------	---

Who	1
-----	---

to	1
----	---

The	1
-----	---

the	1
-----	---

That	1
------	---

Stop	1
------	---

quick	1
-------	---

over	1
------	---

on	1
----	---

Looks	1
-------	---

Look	1
------	---

lazy	1
------	---

jumps	1
-------	---

Jumping	1
---------	---

How	1
-----	---

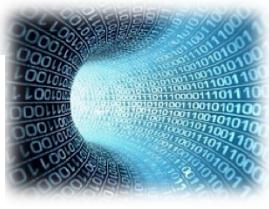
Guests	1
--------	---

Dogs	1
------	---

brown	1
-------	---

Blue	1
------	---

Eyed	1
------	---

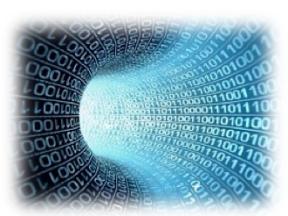


Dog	1
Dog	1
Dog	1

FOX	1
FOX	1
FOX	1

A	1
A	1

Like	1
Like	1



Key Partitioning

Key Value

brown	1
-------	---

dog.	1
------	---

fox	1
-----	---

jumps	1
-------	---

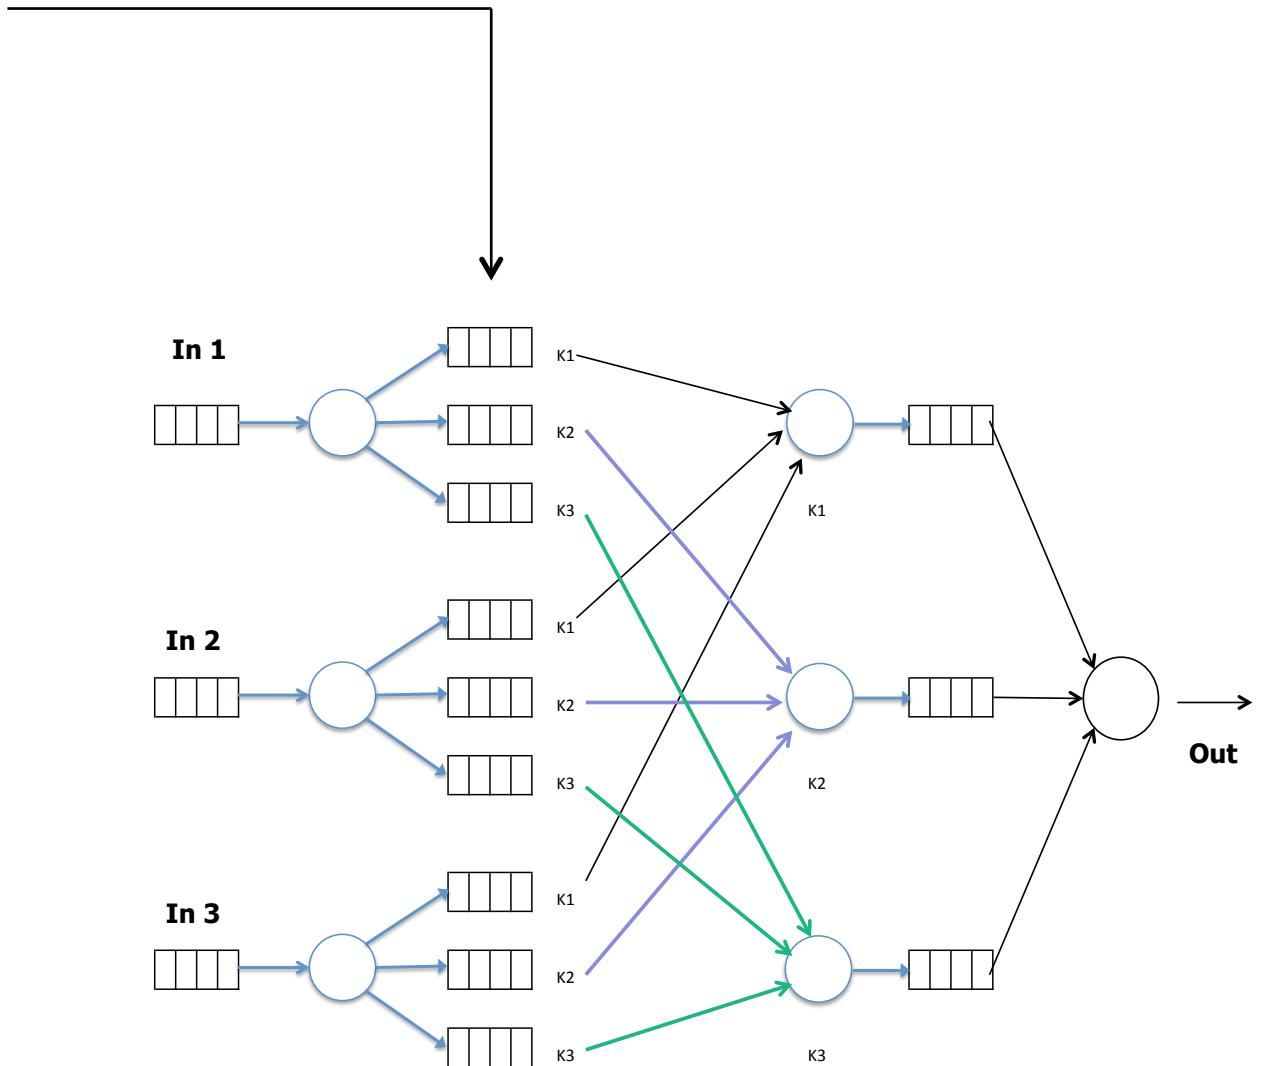
lazy	1
------	---

over	1
------	---

quick	1
-------	---

The	1
-----	---

the	1
-----	---



Key Partitioning



Key Value

brown	1
-------	---

dog.	1
------	---

fox	1
-----	---

jumps	1
-------	---

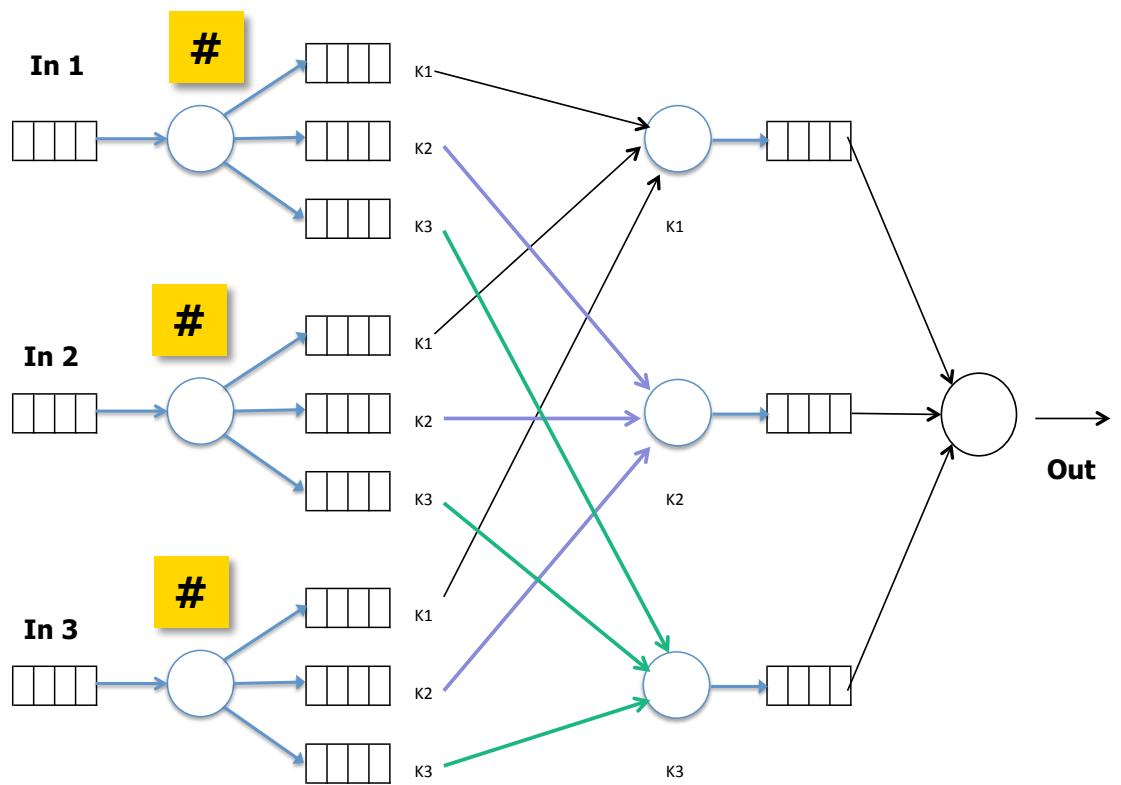
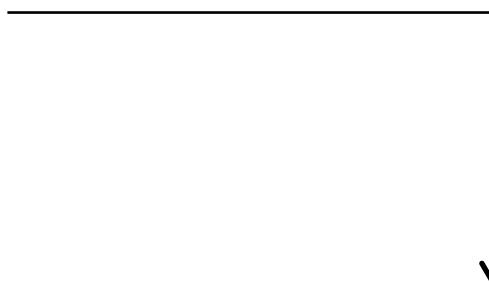
lazy	1
------	---

over	1
------	---

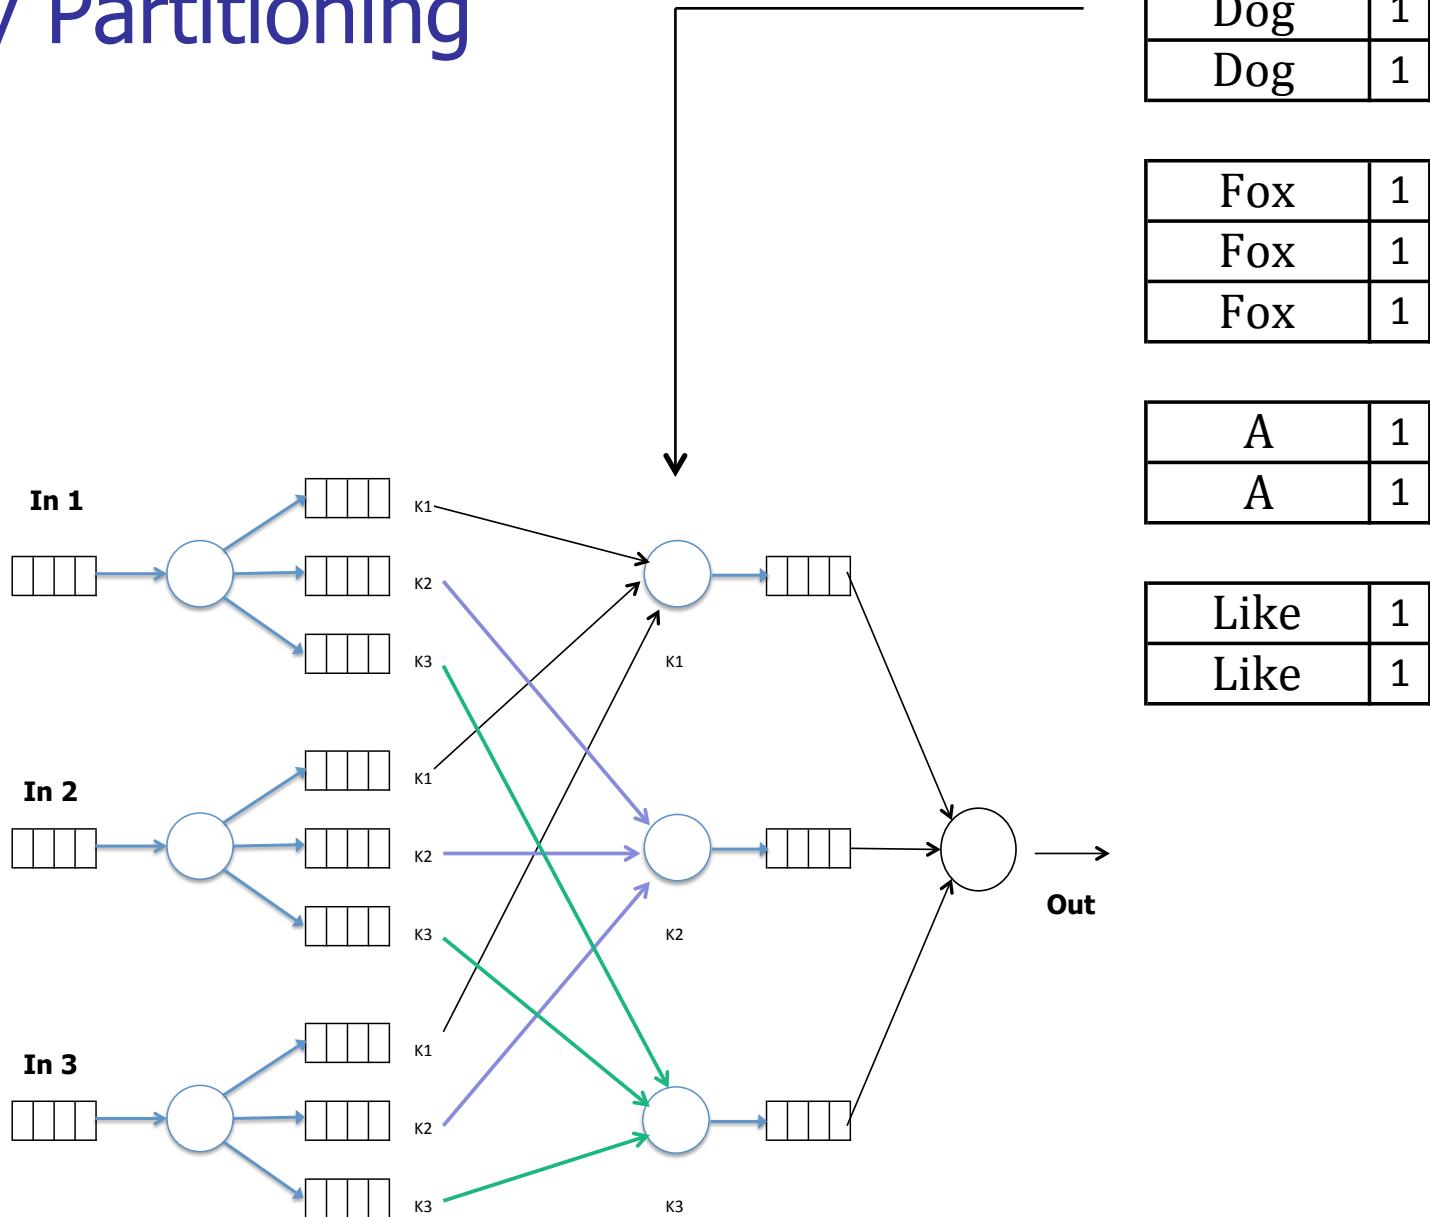
quick	1
-------	---

The	1
-----	---

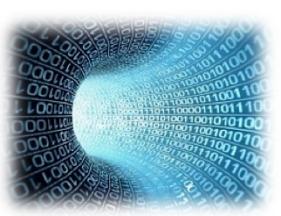
the	1
-----	---



Key Partitioning



Reduce

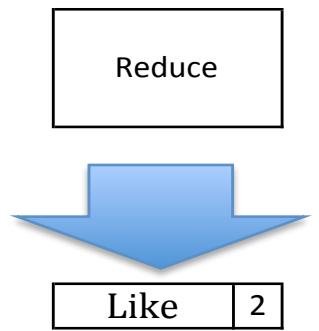
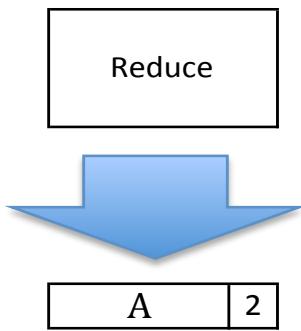
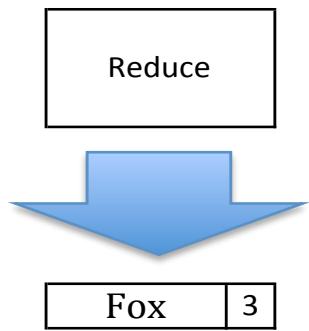
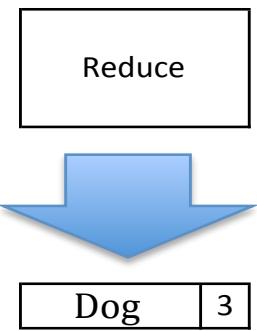


Dog	1
Dog	1
Dog	1

Fox	1
Fox	1
Fox	1

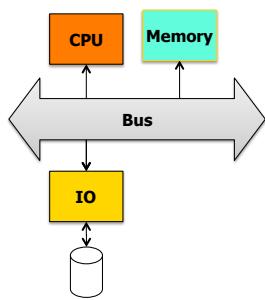
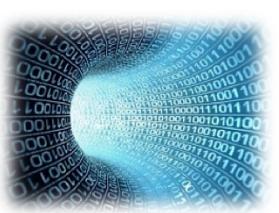
A	1
A	1

Like	1
Like	1

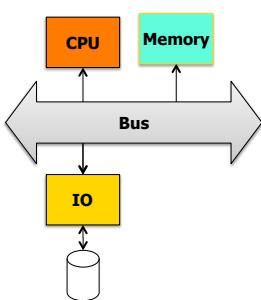


- Reduce tasks take many values for a single key output from map tasks and summarize them into a single output.

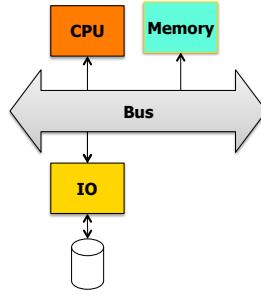
Reduce



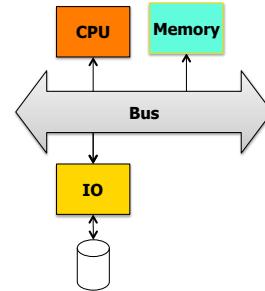
Dog	1
Dog	1
Dog	1



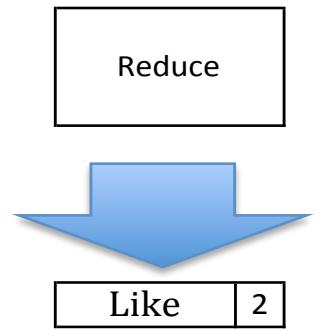
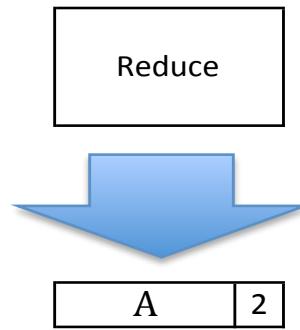
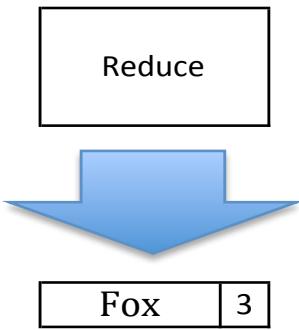
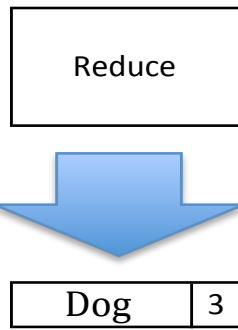
Fox	1
Fox	1
Fox	1



A	1
A	1

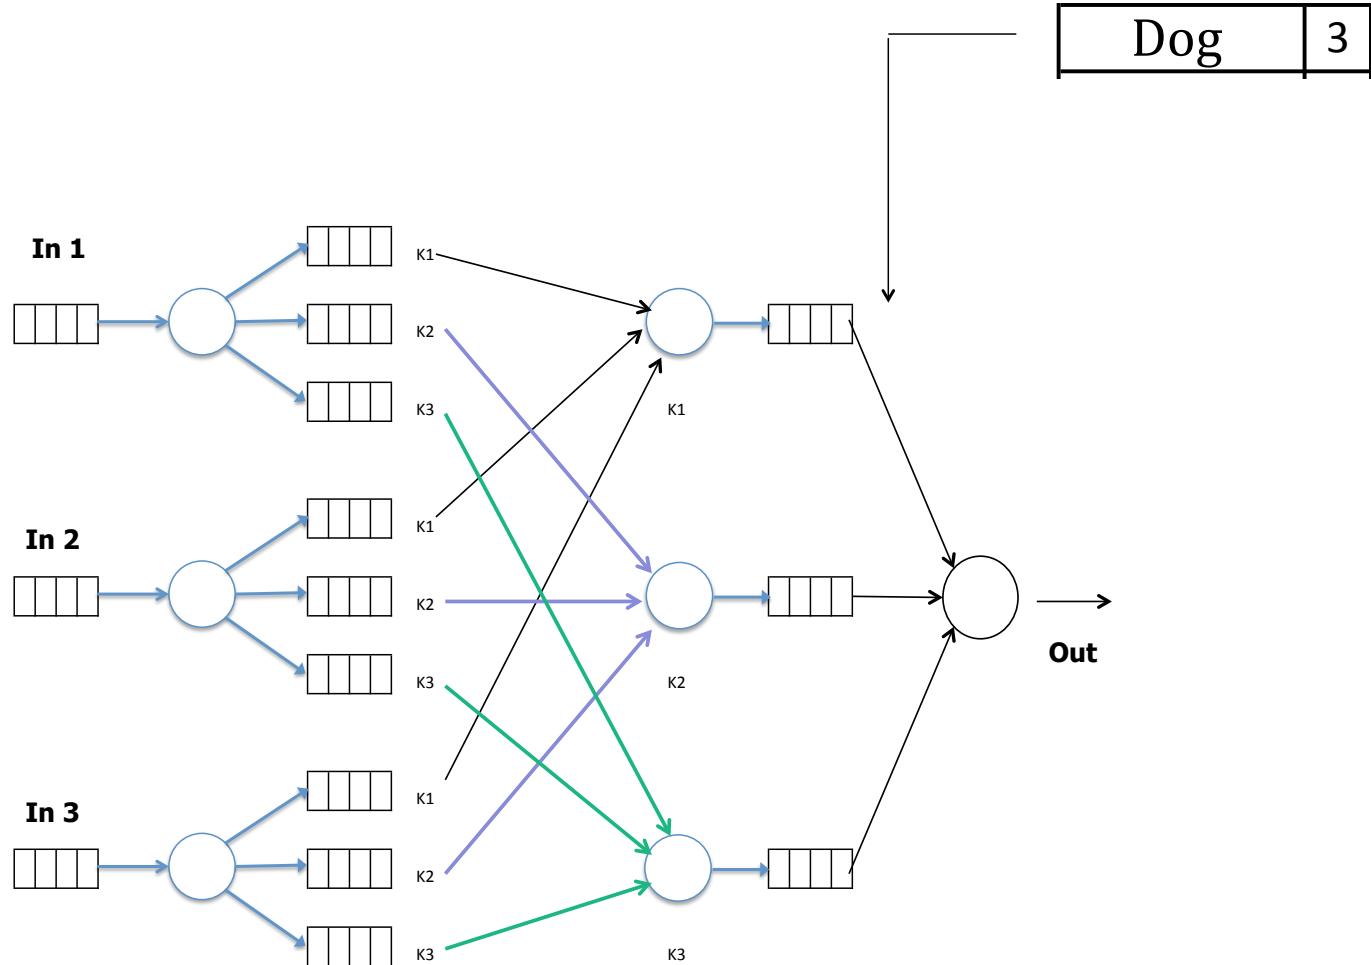
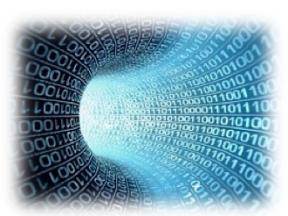


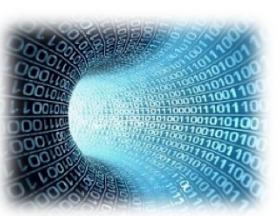
Like	1
Like	1



- Each reducer operates on the result of a single key, so it can be parallelized by key.

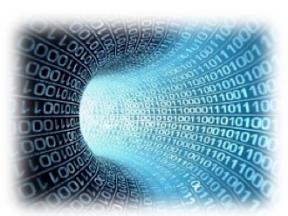
Reduce



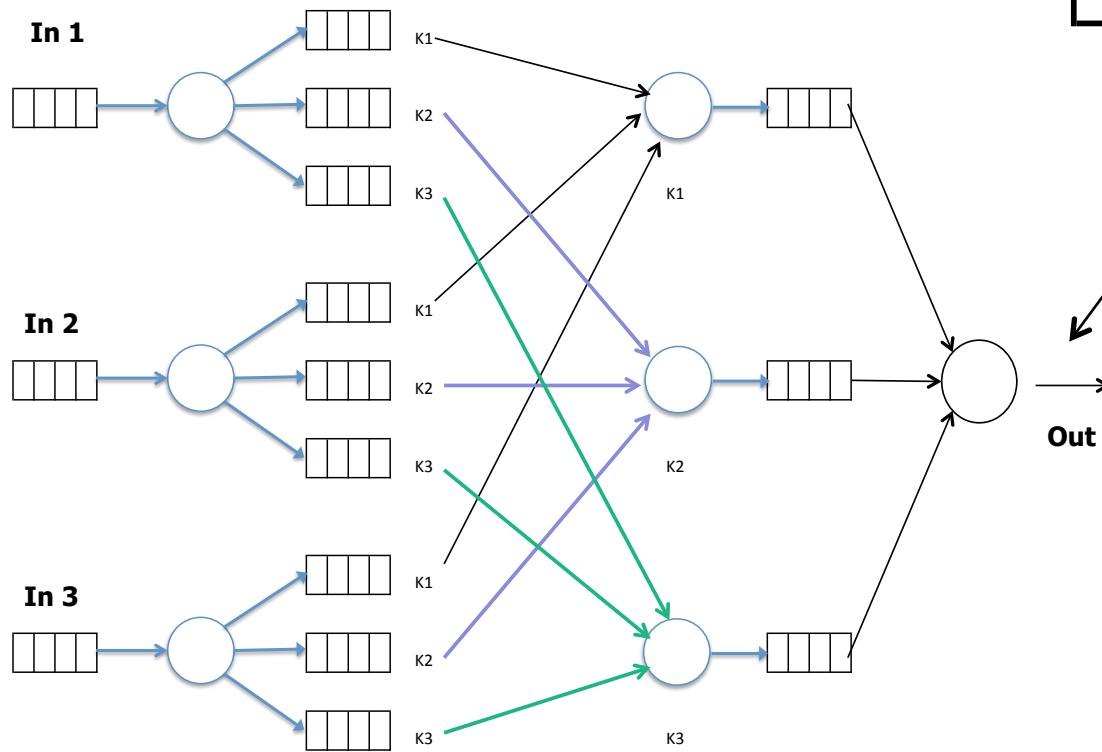


Output

Dog	3
Fox	3
A	2
Like	2

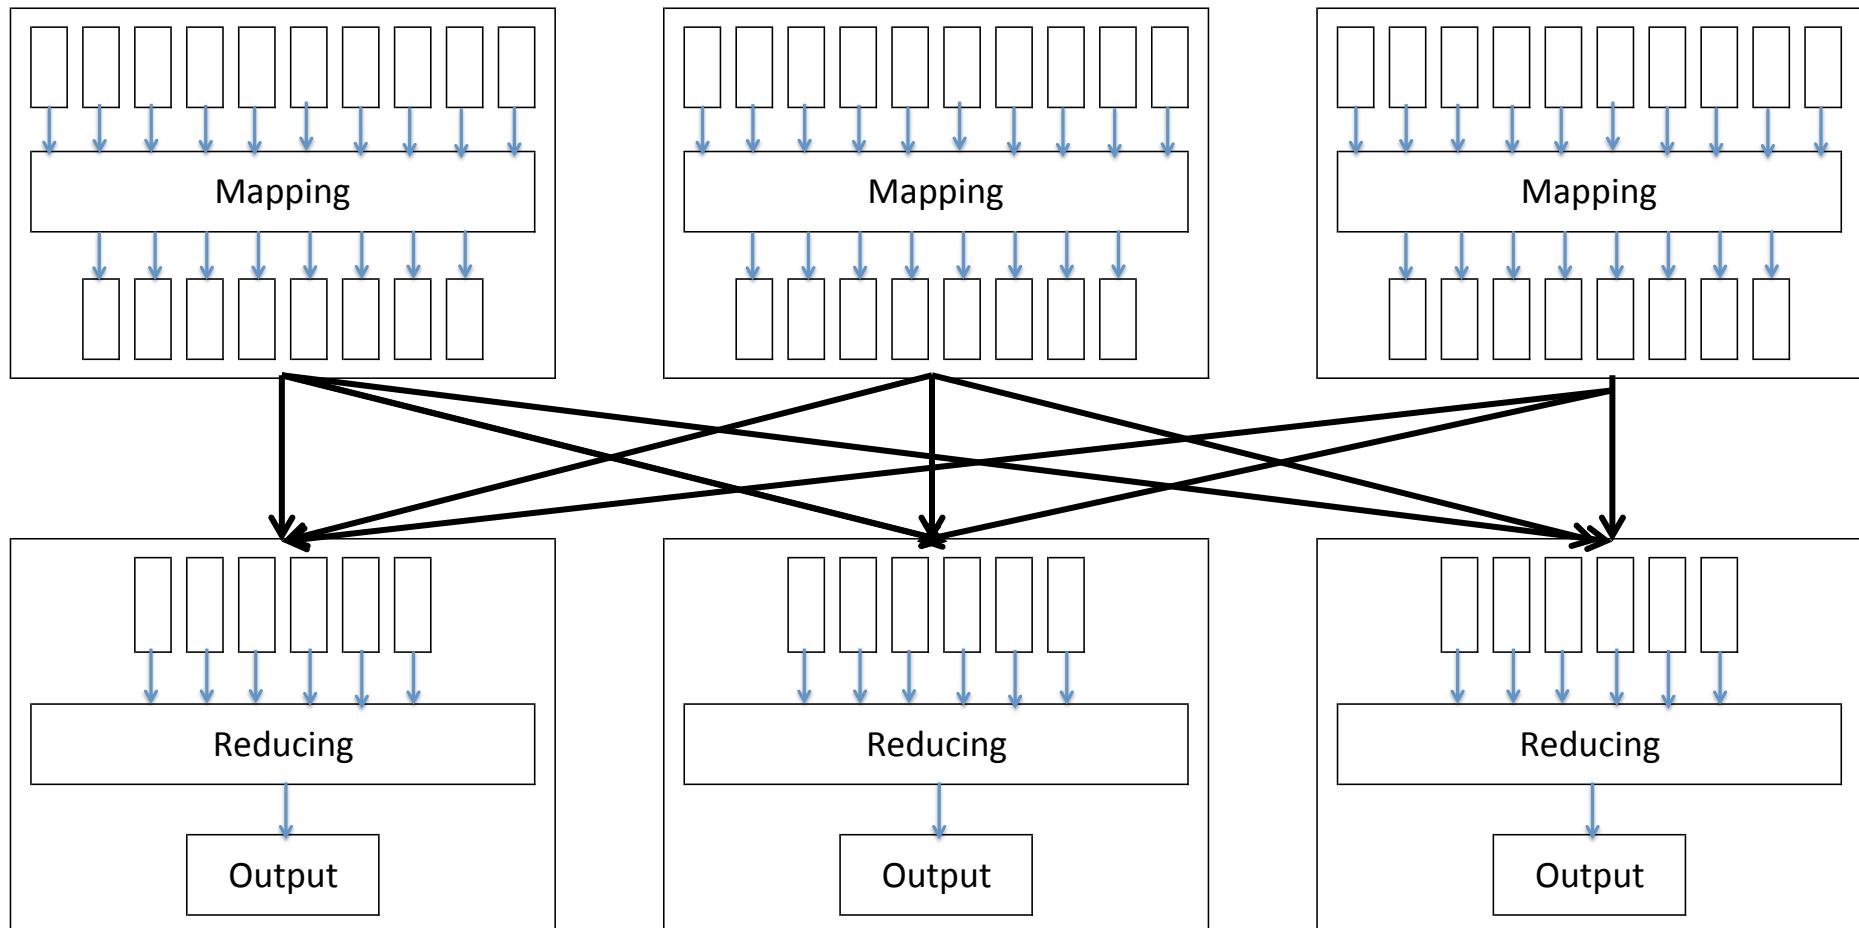
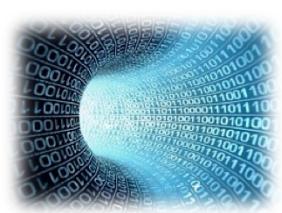


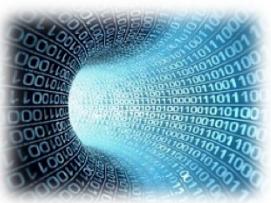
Output



Dog	3
Fox	3
A	2
Like	2

Map Reduce Pipe Line

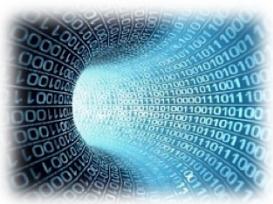




Open source version

- <http://hbase.apache.org/>
- Apache HBase™ is the Hadoop database, a distributed, scalable, big data store.

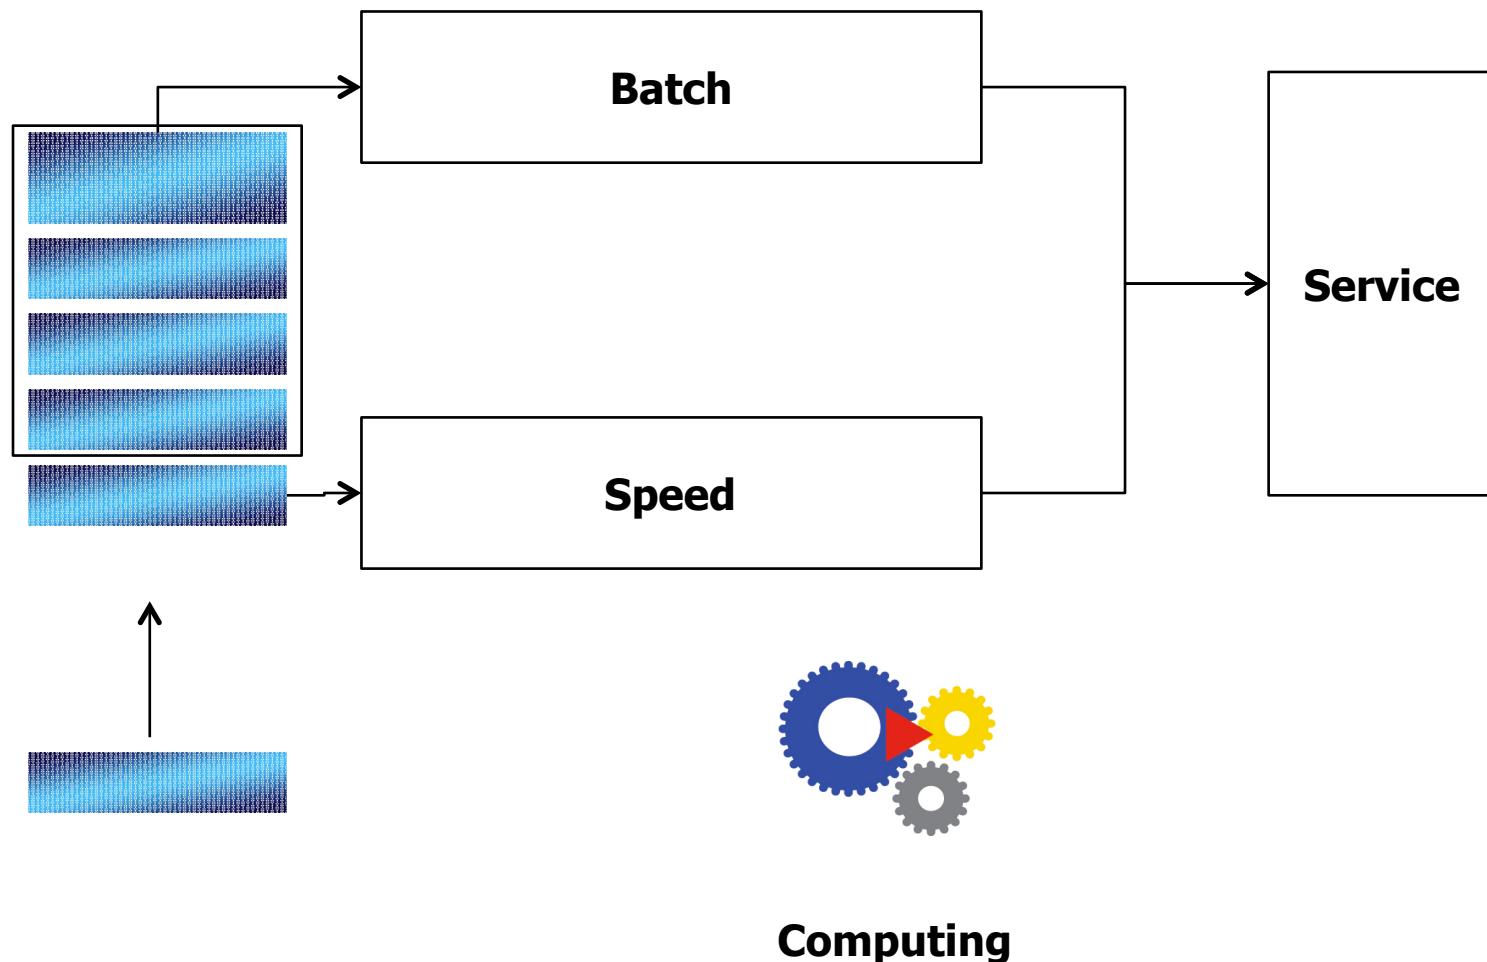
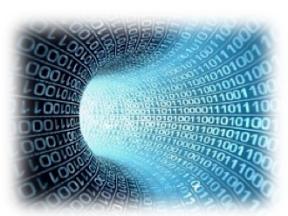


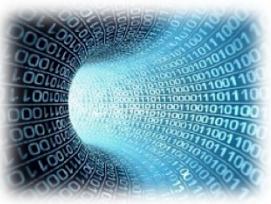


Lambda Architecture

- Lambda architecture provides a combined solution of realtime (speed) data with batch data.
- The Mapreduce Latency for updates is up 2 to 3 hours.
- Intermediary updates are made during the Latency Time

Lambda Architecture





Lambda architecture

- 1) Speed/event layer :
 - Stream processing architecture.
 - The data are appended to the master data periodically over a time window.
 - Real time views
- 2) Batch layer :
 - Manages the master dataset:
 - Immutable, append-only raw data and pre-computing batch views.
 - Batch views
- 3) Serving layer:
 - Merges batch and real-time views.

